

Efficient MLIR Compiler Design: Vectorization for Presburger Library

Zhou Qi



4th Year Project Report
Computer Science
School of Informatics
University of Edinburgh
2023

Abstract

This report presents a faster implementation for the core `pivot` function of MLIR’s presburger library. Its hot loop is element-wise overflow-checked multiplication and addition on an input matrix of low dimension and mostly small value elements.

The current approach of upstream is element-wise multiplication and addition on transprecision integer matrices, from `int64_t` to `LargeInteger`. This can be improved by efficiently utilizing hardware resources, taking advantage of SIMD, and reducing the bit width for every element: the compiler is not capable of automatically generating vectorized instructions for element-wise transprecision computing, and `int64_t` has a much larger bit width than what is typically used for most of the elements in the matrix. Additionally, extra arithmetics are required to perform overflow checking for `int64_t`, resulting in significant performance overhead. This report “innovates” the `int23_t`¹ datatype, a 23-bit integer datatype that utilizes the 23-bit mantissa of a 32-bit floating point, to address these issues. The faster “pivot” performs matrix-wise transprecision computing, targeting 99% **TODO: confirm this number** of the case where elements fit inside `int23_t`. Overflow awareness overhead is almost free, as floating point imprecision implies `int23_t` overflow (See Section ??) and can be captured by a status register. It takes as low as 1 ns to check the status register in the pipeline, and only takes 9 ns to reset the status register. Additionally the status register is only cleared once before a sequence of `pivot` calls, making the average cost of clearing the status register per `pivot` negligible.

On a 30-row by 16-column example matrix, it performs 30 times faster than the upstream scalar implementation. The time cost of a single `pivot` call is reduced from 550 ns to 18.6 ns. **TODO: replace this with actual MLIR benchmark result**

TODO:Question: mention int16 here?

TODO:Reminder: we don’t use int16 for (1) compatibility (2) only slightly faster than float

¹This is not really a innovation. It it a common technique on GPUs because often they are more capable on floating points than integers. See Section 1 for more history and detail.

Research Ethics Approval

This project was planned in accordance with the Informatics Research Ethics policy. It did not involve any aspects that required approval from the Informatics Research Ethics committee.

Declaration

I declare that this thesis was composed by myself, that the work contained herein is my own except where explicitly stated otherwise in the text, and that this work has not been submitted for any other degree or professional qualification except as specified.

(Zhou Qi)

Acknowledgements

Table of Contents

1	Introduction	1
2	Background	4
2.1	Presburger library	4
2.2	Modern CPU micro-architecture	5
2.3	Floating points	6
2.3.1	IEEE 754	6
2.3.2	Fused-multiply-add	6
2.3.3	Representing “int23_t” and “int52_t” using floating points	7
2.4	Google Benchmark	8
2.5	llvm-mca	10
3	Experiments with Toy Example	13
3.1	Vectorization method	13
3.1.1	Clang’s automatic vectorization	13
3.1.2	Clang’s vector datatype and AVX intrinsics	14
3.1.3	Evaluation	15
3.2	Matrix data structure	18
3.3	Matrix element data type	19
3.3.1	Width	19
3.3.2	Overflow checking for integers	20
3.3.3	Overflow checking for floating points	22
3.3.4	Comparing int16_t and float	24
4	Pivot	26
4.1	Implementation and Optimization	26
4.1.1	Matrix-wise transprecision	26
4.1.2	Double buffering	26
4.1.3	Alignment and padding	27
4.1.4	Reduce number of matrix index computation	27
4.1.5	Loop unrolling	27
4.1.6	Vector size specialization	27
5	Conclusion and Future work	28
	Bibliography	29

Chapter 1

Introduction

MLIR, Multi-Level Intermediate Representation, is a infrastructure for building reusable and extensible compilers. Its aim is to reduce fragmentation in domain specific languages and heterogeneous hardware [4]. Its Presburger library provides polyhedral compilation techniques to make dependence analysis and loop optimization [3] and cache modeling [7]. Presburger arithmetics involves determining whether conjunction of linear arithmetic constraints is satisfiable [2], and can be solved using the simplex method of linear programming, with its core function "pivot" consumes **TODO: find the number** % [5] of the runtime.

The `pivot` function involves two multiplication and one addition operation on every element in a matrix. Notably, the input matrices for this library tend to exhibit characteristics of small values and low dimensionality. For example, 90% of test cases work with 16-bit integers that never overflow, and 74% of isl's runtime is spent on test cases that we can compute using 16-bit integers and matrices with at most 32 columns [6]. These properties can be leveraged to take advantage from modern micro-architectural hardware resources, thereby accelerating the process.

Currently, the source code in MLIR upstream adopts a nested for-loop to iterate through every element of the matrix in a transprecision manner. Each number in the matrix can either be `int64_t` or `LargeInteger`. The algorithm starts by using `int64_t`, in case of overflow, it switches to the `LargeInteger` version. This approach is computationally expensive and inefficient, for the following reasons:

1. `int64_t` has a much larger bit width than what is typically used for most of the elements in the matrix,
2. the compiler is not capable of automatically generating vectorized instructions to further optimize the process,
3. overflow is checked manually through additional arithmetic operations.

To propose a faster alternative of the `pivot` function, we could consider constructing a new `pivot` algorithm that satisfies the following conditions:

1. Utilize SIMD: preliminary benchmarks (see Section ??) indicate 8x **TODO:**

verify this number performance improvement on a simple vector element-wise add example.

2. Use small bit width for every element: reducing bit width by half doubles the amount of numbers packed into a single vector register, and essentially reduces the instruction count by half (See Table 3.3).
3. Fast overflow checking: for integers, overflow has to be checked manually and this introduces 60% **(TODO: verify this number?)** overhead, as benchmarks in the Section 3.3.2.1 shown. This is because the x86 architecture does not provide status registers to indicate integer overflow. However, there is one for floating points, making floating points overflow detection almost free.

Previous work [6] have implemented a modified version of `pivot` that utilizes `int16_t` and is designed for matrices with 32 columns or less [6]. This approach offers the advantage of being able to pack a row of 32 elements into a single AVX-512 register and addresses issues 1 and 2. However, overflow is still checked manually, causing 4x or 5x more instruction count (Section 3.3.2.1). Moreover, this approach introduces a new disadvantage, AVX-512 extensions are required for the support for vectorized `int16_t`, this is very rare among CPU manufactured in the last decade (Section 2.2).

An alternative approach is to do 23-bit or 52-bit integer operations using float (32-bit floating point) or double (64-bit floating point) respectively. Though floating points are notorious for precision issues, they are reliable when representing integers that fit inside their mantissa, 23 bits for float and 52 bits for double ¹. When the result of some integer computation exceeds the bit size of the mantissa, floating point imprecision almost always occurs and a status register will be set automatically (Section 3.3.3). Comparing to `int16_t`, even though vector size is sacrificed as there does not exist support for 16-bit floating point `half`, using floating points could still potentially be faster, because overflow checking overhead can be significantly reduced. With floating points, the cost of overflow checking is the time spent on resetting the status register once at the beginning of a sequence of calls to `pivot`, plus reading it in each `pivot` call. Benchmarks **TODO: some figure** indicate that even though reading the register and resetting may take 4.5 ns and 9.4 ns respectively, the effective total overhead is as low as 1 ns. The average cost of resetting per pivot can be treated as negligible, while the superscalar and out-of-order execution pipeline hides the latency of reading the status register with pending memory operations.

This report will first analyze the capability of modern CPU micro-architecture, especially Zen4, through a matrix element-wise fused-multiply-add toy example under the various configurations regarding vectorization methods, matrix data structures, element data types and data widths (Section 3).

It is discovered that optimal performance can be achieved by selecting clang builtin vector type as vectorization methods and use flat list as matrix data structure. However, it is quite difficult to decide whether `int16_t` or `float` is better, because

¹IEEE 754 specification is introduced in Section 2.3.1

the former benefits from bigger vector size and less instruction count, while the latter has minimal overhead on overflow checking.

Then two detached versions of `pivot` function from the Presburger library are built from the most optimal configurations derived from the toy example, one using `int16_t` and the other using `float`. Some further optimizations were made by inspecting `perf` reports and assembly, including:

1. Reduce memory operation
2. Unroll loops
3. **TODO: what are the other optimizations?**

TODO: Update here after Integrating into library

Chapter 2

Background

2.1 Presburger library

The FPL paper obtained 465,460 representative linear problems encountered during integer set coalescing by analyzing linear programs in cache analytical modeling, polyhedral loop optimization, and accelerator code generation. It is found that most of the constraint matrices are low in dimensionality and small in the value of each element. Specifically, more than 99% of the coefficients require less than 10 bits and 95% of them are less than 20 columns [5]. Thus, most of the rows fit inside a 512-bit vector register of 32 `int16_t` elements, and a row operation can be done in a single instruction.

However, in rare and corner cases, there can be larger coefficients up to 127 bits. Practically, the upper bound of coefficient size is unknown, making it required to have arbitrary precision arithmetic `LargeInteger` as a backup. Also, the maximum observed column count is 28 and there is not a certain maximum column count as well.

Therefore the FPL paper presents a 3-way transprecision implementation for the Presburger library, from row-wise vectorized `int16_t` to element-wise scalar `int64_t` and element-wise scalar `LargeInteger`, as illustrated in Figure 2.1. But unfortunately the MLIR upstream only presents a 2-layer transprecision, consisting of element-wise scalar operation using `int64_t` and `LargeInteger`. The `int16_t` version is not merged with the upstream for two reasons:

1. `int16_t` vectors require AVX-512 ISA extension, but hardware support are rare (Section 2.2).
2. Despite the `int16_t` version is fast `TODO: find how much faster in FPL paper`, overflow checking overhead is ??%`TODO: find how much is overhead` [6]. Using floating points could significantly reduce this overhead and potentially be faster (Section 2.3).

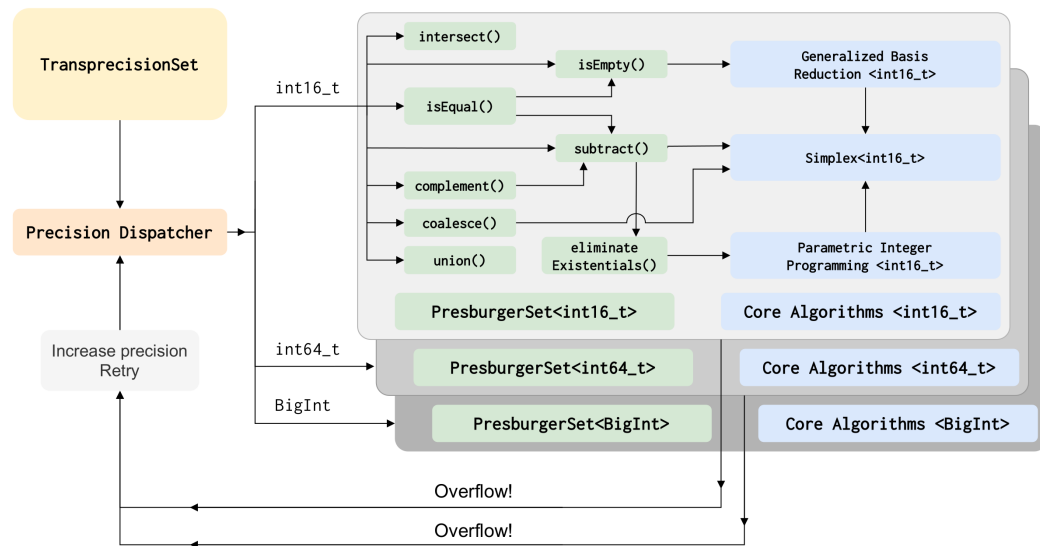


Figure 2.1: The The architecture of FPL.

2.2 Modern CPU micro-architecture

A recent trend in x86-64 architecture’s development is to include AVX-512 instruction set architecture (ISA) extension. AVX-512 succeeds AVX2, the vector width is increased from AVX-2’s 256 bits to 512 bits. AVX-512 also provides new instructions, for example, `int16_t` saturated addition.

Even though its specification was released by Intel in 2013, it had been unpopular, as it did not bring practical performance improvements. The primary reason was that it consumed a lot more power than usual, causing severe overheating. The micro-architecture Skylake from Intel, and its AVX-512 enabled counterpart Skylake-X is a classic example. Skylake provides 2 256 bits FMA AVX-2 execution units¹ and Intel provides 2 512-bit AVX-512 FMA units by fusing the existing AVX-2 units into a AVX-512 unit, then introduces an additional FMA AVX-512 unit [8]. The additional AVX-512 unit increases the heat flux density of the chip, causing server thermal throttling issues.

Intel attempted to mitigate this problem by introducing the “AVX-offset” mode. When a workload involving AVX-512 instructions is encountered, the CPU automatically enters AVX-offset mode and reduces its clock frequency [9]. However, in practice it is more common to have a mix of control flow, SSE and AVX-512 instructions, causing many scenario could run faster if the additional AVX-512 is disabled (and does not thermal throttle) [10].

AMD’s implementation of AVX-512 in Zen4 has slightly less computing power, but much more efficient. Zen4 can be considered as modernized version of Zen3 or Zen2, where Zen2 and Zen3 supports AVX2 by providing 2 FADD units ² and

¹Fused-multiply-add (FMA) execution units are a type of floating point execution units, capable of doing addition, multiplication or both in a single instruction. See Section 2.3.2.

²Floating-point add units (FADD) can execute addition instructions only. They may be

2 FMA units of 256-bit width [11]. Zen4 "double-pumps" these existing circuits to create a single 512-bit FADD and a single 512-bit FMA, without introducing any new execution units for floating point arithmetic [10]. Zen2 and Zen3 are reputable for its high performance per watt [12], and it is expected to be better on zen4 with more advanced lithography. Benchmarks (TODO Some section) indicate that this design indeed does not cause throttling in AVX-512 workloads [10].

Additionally, rebuilding existing software to target AVX-512 may bring slight performance improvement. One benefit of AVX-512 is that it reduces front-end pressure. In the case of Zen4 micro-architecture, though the back-end is possible to commit 2 AVX-2 FADD and 2 AVX-2 FMA every cycle, the front-end has to dispatch 4 instructions per cycle, which is quite difficult. The equivalency in AVX-512 only takes 2 instructions, this is much more likely to be sustained by the frontend [10].

2.3 Floating points

2.3.1 IEEE 754

IEEE 754 is the standard for representing and manipulating floating-point numbers in modern x86 computers. The standard defines several different formats for representing floating point numbers, the most common ones are 32-bit single precision (float) and 64-bit double precision (double). For each format, it specifies how many bits are used to represent the sign, exponent, and mantissa.

For float and double, the sign bit is a single bit that indicates whether the number is positive or negative. As figure 2.2 shows, there are 8 bits and 11 bits for exponent in float and double respectively, to represent represents the order of magnitude. The remaining 23 bits in float and 52 bits in double are mantissae, to store the fractional part of the number. The value of a floating point number can be computed through this formula: $(-1)^s * 2^{(e-B)} * (1 + f)$ where s is sign, e is exponent, f is mantissa and B is a constant bias value: 127 for float, 1023 for double.

2.3.2 Fused-multiply-add

After doing floating point arithmetic, it is required to normalize the result of floating-point arithmetic before it can be used further. However, by feeding the result of a floating-point multiplication (FMUL) directly into the floating-point addition (FADD) logic without the need for normalization and rounding in between, a fused multiply-add (FMA) operation is effectively created: $Y = (A * B) + C$, where A , B and C are the operands, Y is the result [14].

FMA saves cycles and reduces accumulation of rounding errors, while at the same time not adding significant complexity to the circuit. A FMA execution unit is

considered as simplified FMA




Figure 2.2: IEEE 754 single (32 bits) and double (64 bits) precision floating point.
[13]

capable to do FMUL, FADD, and FSUB as well:

Addition: $Y = (A * 1.0) + C$

Multiplication: $Y = (A * B) + 0.0$

Subtraction: $Y = (A * -1.0) + C$

This is a useful feature in many numerical computations that involve simultaneous multiplication and addition operations, such as dot product and matrix multiplication. Since the pivot function performs multiplication and addition between the pivot row, some constant value and each row in the matrix, the performance of FMA is critical to the overall efficiency of the algorithm.

2.3.3 Representing “int23_t” and “int52_t” using floating points

There is a common stereotype that floating-point numbers are unreliable and likely to be imprecise, and are often illustrated in popular memes (as shown in Figure 2.3). However, when storing integer values inside floating points, floating points can be quite reliable.

Specifically, given that the mantissa part of a float consists of 23 bits, inexactness never occur when representing integers less than 23 bits (ignoring the sign bit). Furthermore, in case of of an integer overflow, floating point imprecision almost always occurs and a corresponding status registers is set. The same concept applies to double data types, which have a mantissa consisting of 52 bits.

This mechanism is reliable, as floating-point inexactness always implies integer-inexactness. For a integer value with bit width greater than the mantissa size, floating point rounding is triggered in order to fit most significant bits of the integer in the mantissa, and then adjust the order of magnitude in the exponent accordingly. The lower bits of the mantissa are truncated, therefore causing imprecision.




Figure 2.3: A floating point meme: $0.1 + 0.2 = 0.30000000000000004$

In some rare cases, integers longer than the size of mantissa can be represented in floating points precisely. An examples of such numbers are large powers of 2, like $2^{30} = 0x40000000$. Its binary representation in float is:

Sign : 0

Exponent: 10011011

Mantissa: 000000000000000000000000

Despite being greater than the size of the mantissa, they are normalized rather than being rounded, and therefore does not break the mechanism of representing integer in floating-points.

Throughout history, floating point processing units in GPUs have been utilized for fast integer arithmetic, for example, modular exponentiation [15] and RSA algorithms [16], because often the architecture of GPU prioritizes floating points performance rather than integers. **TODO: this paragraph here? or introduction?**

2.4 Google Benchmark

Google benchmark is a library to measure the performance of a code snippet. It provides unit-test like interfaces to setup benchmarks around a code snippet [1]. The given example from <https://github.com/google/benchmark> is self-explanatory for its usage:

```
#include <benchmark/benchmark.h>
```

```
static void BM_SomeFunction(benchmark::State& state) {
    // Perform setup here
    for (auto _ : state) {
        // This code gets timed
        SomeFunction();
    }
}

// Register the function as a benchmark
```

```
BENCHMARK(BM_SomeFunction);
// Run the benchmark
BENCHMARK_MAIN();
```

The library first starts a timer, repeatedly executes its core loop: `for (auto _ : state) ...` multiple times then pauses the timer. This method ensures that the results are consistent and minimizes the overhead required for recording the timing information.

Executing the benchmarks will not only report both elapsed real time and CPU time, but also much other useful information to help reduce variance.

```
Running ./build/example
***WARNING*** CPU scaling is enabled, the benchmark real time
measurements may be noisy and will incur extra overhead.
Run on (32 X 5800.00 MHz CPU s)
CPU Caches:
  L1 Data 32 KiB (x16)
  L1 Instruction 32 KiB (x16)
  L2 Unified 1024 KiB (x16)
  L3 Unified 32768 KiB (x2)
Load Average: 8.10, 5.14, 1.14
```

Benchmark	Time	CPU	Iterations
BM_SomeFunction	18.5 ns	18.5 ns	37935734

The warning: “CPU scaling is enabled, the benchmark real-time measurements may be noisy and will incur extra overhead.” is saying that CPU clock frequency is not consistent. It can be dynamically determined by the governor algorithm, according on the system’s needs. For example, with the **performance** governor, the OS locks the CPU to the highest possible clock frequency, specified at `/sys/devices/system/cpu/cpu*/cpufreq/scaling_max_freq`, while the **ondemand** governor will push the CPU to the highest frequency on demand and then gradually reduce the frequency as the idle time increases [17].

However, it is also dependent on the manufacture and other hardware constraints. By default, both Intel (Turbo Boost) and AMD (Precision Boost Overdrive) have support for raising clock frequency, beyond the control of the governor [18]. On the other hand, CPUs have self-protecting thermal throttling mechanisms that reduces its clock frequency and voltage when it is too hot.

The benchmark mentioned in this report were performed on a AMD 7950x desktop computer. The computer system went through the following these setups for consistent results:

1. Set the governor to **performance**,
2. Disable AMD Precision Boost Overdrive (or Intel Turbo Boost),
3. Lock clock frequency at a 5 GHz, or any desired fixed value,
4. Make sure heat dissipation is working properly.

2.5 llvm-mca

`llvm-mca`, LLVM Machine Code Analyzer, is a tool to analyze performance of executing some instructions on a specific CPU micro-architecture, according to scheduling information provided by LLVM [19].

By supplying `llvm-mca` with a piece of assembly code and the target micro-architecture codename, `llvm-mca` reports various metrics to indicate how fast the given instructions will execute on the specified micro-architecture. It first summarizes the instruction per clock (IPC) and throughput of the entire instruction block, then gives detailed information about each instruction, including number of uOps, latency, throughput, potential load, store and side effects. `llvm-mca` also reports resource pressure in terms of arithmetic units and memory load or store units. When the optional `-timeline` flag is prompted, `llvm-mca` illustrates a timeline view of the analyzed code, showing how instructions progress through the pipeline stages of the target processor. The timeline helps understand the capability of complicated out-of-order superscalar architecture.

An example is provided below. The analysis from `llvm-mca` indicates that a `znver2` (Zen2) CPU can repeatedly execute a combination of `vmovaps`³ and `vfmadd213ps`⁴ instructions every 1.3 cycle. This translates to 2.74 instructions per cycle. The output from `llvm-mca` is slightly modified and truncated to fit into limited page size.

```
$ llvm-mca-15 -timeline -mcpu=znver2 ./x.s
```

```
Iterations:      100
Instructions:     400
Total Cycles:    146
Total uOps:      400
```

```
Dispatch Width:  4
uOps Per Cycle:  2.74
IPC:              2.74
Block RThroughput: 1.3
```

Instruction Info:

```
[1]: #uOps
[2]: Latency
[3]: RThroughput
[4]: MayLoad
[5]: MayStore
[6]: HasSideEffects (U)
```

```
[1] [2] [3] [4] [5] [6] Instructions:
1   8  0.33 *          vmovaps (%rdx,%rsi,4), %ymm0
```

³`vmovaps` can be either vector float load or store instruction, depending on how operands are structured.

⁴`vfmadd213ps` is the instruction for fused-multiply-add.

```

1   8  0.33  *           vmovaps  (%rcx,%rax,4), %ymm1
1  12 0.50  *           vfmadd213ps (%r8,%rdi,4), %ymm0, %ymm1
1   1  0.33      *       vmovaps  %ymm1, (%r9,%rax,4)

```

Resources:

```

[0]   - Zn2AGU0
[1]   - Zn2AGU1
[2]   - Zn2AGU2
...
[8]   - Zn2FPU0
[9]   - Zn2FPU1
[10]  - Zn2FPU2
[11]  - Zn2FPU3
[12]  - Zn2Multiplier

```

Resource pressure per iteration:

```

[0] [1] [2] [3] [4] [5] [6] [7] [8] [9] [10] [11] [12]
1.33 1.33 1.34 - - - - - 0.50 - - 0.50 -

```

Resource pressure by instruction:

```

[0] [1] [2] ... [8] [11] Instructions:
0.01 0.38 0.61 ... - - vmovaps  (%rdx,%rsi,4), %ymm0
0.23 0.68 0.09 ... - - vmovaps  (%rcx,%rax,4), %ymm1
0.27 0.24 0.49 ... 0.50 0.50 vfmadd213ps (%r8,%rdi,4), %ymm0, %ymm1
0.82 0.03 0.15 ... - - vmovaps  %ymm1, (%r9,%rax,4)

```

Timeline view:

```

                                0123456789
Index 0123456789                012345

[0,0] D=eeeeeeeeER . . . vmovaps  (%rdx,%rsi,4), %ymm0
[0,1] D=eeeeeeeeER . . . vmovaps  (%rcx,%rax,4), %ymm1
[0,2] D=eeeeeeeeeeeeER . . vfmadd213ps (%r8,%rdi,4), %ymm0, %ymm1
[0,3] D=====eER . . vmovaps  %ymm1, (%r9,%rax,4)
[1,0] .DeeeeeeeeE-----R . . vmovaps  (%rdx,%rsi,4), %ymm0
[1,1] .DeeeeeeeeE-----R . . vmovaps  (%rcx,%rax,4), %ymm1
[1,2] .D=eeeeeeeeeeeeER . . vfmadd213ps (%r8,%rdi,4), %ymm0, %ymm1
[1,3] .D=====eER . . vmovaps  %ymm1, (%r9,%rax,4)
[2,0] . DeeeeeeeeE-----R . . vmovaps  (%rdx,%rsi,4), %ymm0
[2,1] . DeeeeeeeeE-----R . . vmovaps  (%rcx,%rax,4), %ymm1
...

```

However, when evaluating the identical assembly code on a more advanced micro-architecture **znver3** (Zen3), **llvm-mca** reveals reduction in IPC and throughput. This appears to be contradictory with both theoretical expectations and actual benchmarks. After submitting an issue [20], **llvm** maintainers explained that **llvm**'s

scheduling information are hand-crafted using `llvm-exegesis`, a micro-benchmark tool. The issue was subsequently resolved after rerunning `llvm-exegesis` and confirming that `znver3` indeed has higher throughput as expected.

TODO: should I write a paragraph to emphasise that `llvm-mca` is unreliable?

Chapter 3

Experiments with Toy Example

The pivot function does multiply and add for each row in the matrix, therefore the performance of FMA a simple vector toy can be an effective indicator. This chapter reports performance analysis on simple toy examples that do vector add or vector FMA with various setups, including:

1. Vectorization method
 - (a) Clang's automatic vectorization from scalar source code
 - (b) Clang builtin vector datatype with occasional AVX intrinsics
2. Matrix data structure
 - (a) Nested list
 - (b) Flat list
3. Element data width
 - (a) 16 bits: `int16_t`
 - (b) 32 bits: `int32_t`, `float`
 - (c) 64 bits: `int64_t`, `double`
4. Element data type
 - (a) Integer
 - (b) Floating point

3.1 Vectorization method

3.1.1 Clang's automatic vectorization

Clang is capable of generating vectorized instructions from scalar source code, using the flags `-O3 -march=native` on a platform with vector ISA enabled. Starting with an example (Listing 3.1), the simple `vec_add` function adds every element from two arrays and saves it to the third.

TODO: these hex code are actually incorrect, I assume this does not really matter?

Source code

```
#define size 128
void vec_add(float* src1_ptr, float* src2_ptr, float* dst_ptr) {
    for (uint32_t i = 0; i < size; i += 1 ){
        dst_ptr[i] = src1_ptr[i] + src2_ptr[i];
    }
}
```

Assembly snippet of the hot loop, compiled with `-O3 -march=native`
(vectorization on)

```
1458: c4 c1 7c 58 84 87 20    vaddps -0x1e0(%r15,%rax,4),%zmm0,%zmm0
145f: fe ff ff
1462: c4 c1 74 58 8c 87 40    vaddps -0x1a0(%r15,%rax,4),%zmm1,%zmm1
1469: fe ff ff
```

Assembly snippet of the hot loop, compiled with `-O3 -march=native`
`-mno-avx -mno-sse` (vectorization off)

```
120d: d8 44 82 04    fadds  0x4(%rdx,%rax,4)
1211: d9 5c 81 04    fstps  0x4(%rcx,%rax,4)
1215: d9 44 86 08    flds   0x8(%rsi,%rax,4)
```

Listing 3.1

After compiling on a AVX-512 enabled computer and disassembling the binary, it is observed that clang automatically packs 16 float (512 bits) as a operand of the VADDPS instruction.

Alternatively, vectorization could be disabled by adding the `-mno-avx -mno-sse` flags on top of `-O3 -march=native`. These two sets of flags guarantee that the binary are going to be equally optimized, with the only difference been whether vector instructions are generated or not. In this case, scalar instructions `fadds`, `fstps` and `flds` are selected.

3.1.2 Clang's vector datatype and AVX intrinsics

Another approach is to write source code with vectorization in mind in the first place. Clang provides extension that allows programmers to declare a new type that represents a vector of elements of the same data type. The syntax is `typedef ty vec_ty __attribute__((ext_vector_type(vec_width)))`, where `vec_ty` is the name of vector type being defined, `vec_width` is its size and `ty` is the type of the elements in the vector. For example, `typedef int16_t int16x32 __attribute__((ext_vector_type(32)))` defines a 512-bit vector type of `int16x32`, consisting of 32 `int16_t` and fits inside an AVX-512 ZMM register.

After defining a vector datatype, a vector variable can be created by casting from a pointer of the target array. Then arithmetic operators can be applied between the vectors to performed element-wise operations. The previous `vec_add` example can be rewritten as the code snippet shown in Listing 3.2:

Source code

```
#define size 128
#define FloatZmmSize 16
typedef float floatZmm __attribute__((ext_vector_type(FloatZmmSize)));
void vec_add(float* src1_ptr, float* src2_ptr, float* dst_ptr) {
    for (uint32_t i = 0; i < size; i += FloatZmmSize ){
        floatZmm src1Vec = *(floatZmm *)(src1_ptr + i);
        floatZmm src2Vec = *(floatZmm*)(src2_ptr + i);
        floatZmm resultVec = src1Vec + src2Vec;
        *(floatZmm *)(dst_ptr + i) = resultVec;
    }
}
```

Listing 3.2

3.1.3 Evaluation

When comparing the performance of code written with and without the vector type and examining their assembly, it has been discovered that the the automatic vectorization feature in clang can be unpredictable and may lead to undesired behaviors. It operates as a black box and may take a lot of effort to understand its mechanisms. One of the issues is that clang may select a suboptimal vector width.

Consider the `vec_fma` function in Listing 3.3 and 3.4, a slightly more complicated version of the previous `vec_add` example, where an additional array is introduced and the element-wise operation is changed from addition to FMA. The disassembly reveals that clang decides to use FMA vector instructions of 128-bit width, but when vector size is constrained to 512-bit width by defining a vector type, more optimal binary can be generated. Benchmark indicates that the 512-bit vector width version is xx% faster. **TODO: determine this number**

Source code

```

void vec_fma(float* src1_ptr, float* src2_ptr,
            float* src3_ptr, float* dst_ptr) {
    for (uint32_t i = 0; i < size; i += 1 ){
        dst_ptr[i] = src1_ptr[i] * src2_ptr[i] + src3_ptr[i];
    }
}

```

Assembly snippet of the hot loop

```

80e4: c5 fa 10 04 b2    vmovss (%rdx,%rsi,4),%xmm0
80f1: c5 fa 10 0c 81    vmovss (%rcx,%rax,4),%xmm1
80fe: c4 c2 79 a9 0c b8 vfmadd213ss (%r8,%rdi,4),%xmm0,%xmm1
810c: c4 c1 7a 11 0c 81 vmovss %xmm1, (%r9,%rax,4)

```

Listing 3.3

Source code

```

#define size 128
#define FloatZmmSize 16
typedef float floatZmm __attribute__((ext_vector_type(FloatZmmSize)));
void vec_fma(float* src1_ptr, float* src2_ptr,
            float* src3_ptr, float* dst_ptr) {
    for (uint32_t i = 0; i < size; i += 1 ){
        floatZmm src1Vec = *(floatZmm *)(src1_ptr + i);
        floatZmm src2Vec = *(floatZmm *)(src2_ptr + i);
        floatZmm src3Vec = *(floatZmm *)(src3_ptr + i);
        *(floatZmm *)(dst_ptr + i) = src1Vec * src2Vec + src3Vec;
    }
}

```

Assembly snippet of the hot loop

```

82c0: 62 b1 7c 48 10 04 80 vmovups (%rax,%r8,4),%zmm0
82c7: 62 b1 7c 48 10 0c 81 vmovups (%rcx,%r8,4),%zmm1
82ce: 62 b2 7d 48 a8 0c 86 vfmadd213ps (%rsi,%r8,4),%zmm0,%zmm1
82d5: 62 b1 7c 48 11 0c 87 vmovups %zmm1, (%rdi,%r8,4)

```

Listing 3.4

In some cases clang could be even worse, it may fail to recognize vectorization patterns from element-wise loop operations, leading to more reduction in performance. In the `vec_fma` example, by changing the type signature from `float` to `int`, clang decides to dispatch scalar instructions for addition (`add`) and multiplication (`imul`)

completely (Listing 3.5). Their vectorized equivalency `vpadd` and `vpmulld` are more performant options. (Listing 3.6).

Source code

```
#define size 128
void vec_fma(int* src1_ptr, int* src2_ptr,
             int* src3_ptr, int* dst_ptr) {
    for (uint32_t i = 0; i < size; i += 1 ){
        dst_ptr[i] = src1_ptr[i] * src2_ptr[i] + src3_ptr[i];
    }
}
```

Assembly snippet of the hot loop

```
88b0: 44 8b 69 1c  mov    0x1c(%rcx),%r13d
88b4: 44 8b 62 1c  mov    0x1c(%rdx),%r12d
88b8: 45 0f af ee  imul    %r14d,%r13d
88bc: 45 0f af e6  imul    %r14d,%r12d
88c0: 45 01 fd      add     %r15d,%r13d
88c3: 45 01 fc      add     %r15d,%r12d
```

Listing 3.5

Source code

```
#define size 128
#define IntZmmSize 16
typedef int intZmm __attribute__((ext_vector_type(IntZmmSize)));
void vec_fma(int* src1_ptr, int* src2_ptr,
             int* src3_ptr, int* dst_ptr) {
    for (uint32_t i = 0; i < size; i += 1 ){
        intZmm src1Vec = *(intZmm *)(src1_ptr + i);
        intZmm src2Vec = *(intZmm *)(src2_ptr + i);
        intZmm src3Vec = *(intZmm *)(src3_ptr + i);
        *(intZmm *)(dst_ptr + i) = src1Vec * src2Vec + src3Vec;
    }
}
```

Assembly snippet of the hot loop

```
8880: 62 b1 fe 48 6f 04 81  vmovdqu64 (%rcx,%r8,4),%zmm0
8887: 62 b2 7d 48 40 04 80  vpmulld (%rax,%r8,4),%zmm0,%zmm0
888e: 62 b1 7d 48 fe 04 86  vpadd (%rsi,%r8,4),%zmm0,%zmm0
8895: 62 b1 fe 48 7f 04 87  vmovdqu64 %zmm0, (%rdi,%r8,4)
```

Listing 3.6

3.2 Matrix data structure

The most intuitive data structure of a matrix is a list of lists, where each list represents a row and a list of rows is a matrix. In C++ this can be represented using `std::vector<std::vector<T>>`, where `T` could be `float`, `double`, `int32_t`, etc. The `std::vector` class provides an intuitive interface for accessing and modifying elements, making it easy to write code with.

One potential drawback of nested `std::vector` is that it requires two indexing operations to access an element. An alternative implementation is to “flatten” a matrix into a single `std::vector`, by simply concatenating one row after another. To access a specific element, an index can be computed manually using the given row and column: `column_count * row + column`. This reduces half of the memory indexing operation at the cost of additional arithmetic. The differences between the two patterns are illustrated by a example provided in Table 3.7.

	Nested	Flat
Type	<code>std::vector< std::vector<int32_t>></code>	<code>std::vector<int32_t></code>
Structure in Memory	vector of 4 = { vector of 4 = {0, 0, 0, 0}, vector of 4 = {0, 0, 0, 0}, vector of 4 = {0, 0, 0, 1}, vector of 4 = {0, 0, 0, 0} }	vector of 16 = { 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0 }
Accessing row 2, column 3	Index row first: <code>std::vector<int32_t>[2]</code> then index column: <code>std::vector<int32_t>[2][3]</code>	Compute <code>i =</code> <code>col_count * row + col</code> <code>= 4 * 2 + 3 = 11,</code> then index once: <code>std::vector<int32_t>[11]</code>

Table 3.7: This is an example of a 4 by 4 matrix, to highlight the differences between the two matrix data structures: nested list and flat list.

Empirically indexing costs more time than integer multiplication and addition, thereby improving performance. Both the toy example and the pivot function perform sequential load-compute-store operations on each row and each column, allowing the index of the next element to be computed by simply adding the step size or column size, further reducing memory overhead. In fact, the pivot function can be optimized to only compute index once (See Section 4.1.4) by placing the pivot row as first row. Benchmark (Figure 3.1) on the toy example confirms that when there are 16 rows, the nested vector matrix is about 8 ns faster than the flat matrix.




Figure 3.1: TODO: write caption

3.3 Matrix element data type

3.3.1 Width

Since the numbers stored in the matrix are almost always less than 10 bits, using shorter data types can be more advantageous than longer ones because they allow more numbers to be packed into a single vector register (Table 3.3). The number of instructions can be cut by half when data width is reduced to half, and less instruction count always leads to less execution time. Given that the Zen4 micro-architecture provides approximately same amount of execution units for both integers and floating points, it is reasonable to estimate that the execution time is inversely proportional to the bit width of data type. As confirmed in Plot 3.2, `int32_t` and `float` costs nearly same amount of time, while `int32_t` and `double` costs double the amount of time than `int16_t` and `float` respectively.




Figure 3.2: **TODO: write caption**

TODO: fix Table number

	f32	f64	i16	i32	i64
Execution units	1 512-bit FADD + 1 512-bit FMA		2 512-bit ALU		
Fused-multiply-add	Yes		No		only on lower 52 bits, no overflow exception
Saturated add	N/A		Yes: <code>_mm512_adds_epi32</code>	No	No
Multiply higher bits	N/A		Yes: <code>_mm512_mulhi_epi16</code>	No	No
SIMD Floating-Point Exceptions	Overflow, Underflow, Invalid, Precision, Denormal		No		
512-bit vector size	16	8	32	16	8
Overflow checking and time cost	single time overhead: <code>fetestexcept_mxcsr</code> 4.09 ns <code>feclearexcept_mxcsr</code> 9.42 ns		additional arithmetic about 4x more instructions (saturated add or multiply higher bits)		Require more instructions than int16 due to lack of saturated add and multiply higher bits operations.

Table 3.3: A summary for features provided by the Zen4 micro-architecture for different data types [10] [11].

3.3.2 Overflow checking for integers

The x86-64 micro-architecture provides the `seto` instruction to set some byte to 1 if overflow occurred as a result of integer arithmetic. However `seto` only works for scalar operations, there does not exist instruction or status register to indicate whether a previous vector add or multiply instruction produces overflown results or not. Therefore, overflow has to be checked manually by some additional

vector instructions, this would slow down the computation to some extent. Or alternatively or arithmetics have to be carried out on each element individually in a scalar manner, resulting in even worse performance.

One advantage of `int16_t` is that it can be used with AVX-512's saturated add and multiply higher bits vector instructions (Table 3.3), making it possible and convenient to write vectorized and overflow-aware code. However, Zen4's implementation of AVX-512 extension does not provide equivalent instruction for `int32_t` or `int64_t`, and therefore must be processed as scalar values.

3.3.2.1 Implementation of vectorized `int16_t` overflow checking

By comparing the result of a conventional addition and saturated addition, it indicates whether an addition has gone overflown or not. In case of overflow, with saturated add, the result always retains at the maximum possible value of `int16_t`: `0x7FFF`, while the result of a conventional add is always smaller, because the overflowed output from conventional addition can't go all the way around and become `INT16_MAX` again. In the two's complement binary form for integer, the overflow sum is "trapped" in the negative number space. For example:

```
INT16_MAX + 1 = INT16_MIN = -32768,
INT16_MAX + 2 = -32767,
...
INT16_MAX + INT16_MAX = -2,
```

For multiplication, as two 16-bit numbers produces 32-bit products but only lower 16 bits can be stored, overflow can be detected by checking whether any of the upper 16 bits are set.

Inspecting these approaches from a instruction-level perspective (Table 3.8), when overflow is ignored, both add and multiply takes 1 instruction, `vpaddw`¹ and `vpnullw`². To obtain and process overflow-related information, an additional computation instruction `vpaddsw`³ or `vpmulhw`⁴ is required, followed with 2 or 3 comparison, shuffling and branch instructions: `vpsraw`⁵, `vpcmpneqw`⁶ and By enabling `kord`⁷. overflow checking, it brings 4x or 5x more instruction count and 60% more runtime [6]. (TODO: verify this number?)

3.3.2.2 Implementation of scalar `int32_t` and `int64_t` overflow checking

Clang's language extension provides functions to perform overflow-checked integer arithmetics:

-
- ¹Vector add for `int16_t`
 - ²Vector multiply lower half bits for `int16_t`
 - ³Vector saturated add for `int16_t`
 - ⁴Vector multiple higher half bits for `int16_t`
 - ⁵Shift packed data right arithmetic
 - ⁶Compare packed data for equal
 - ⁷Bitwise logical OR masks

	Addition	Multiplication
Disabled	<code>vpaddw %zmm4,%zmm2,%zmm3</code>	<code>vpmullw %zmm1,%zmm3,%zmm2</code>
Enabled	<code>vpaddw %zmm4,%zmm2,%zmm3</code>	<code>vpmullw %zmm1,%zmm3,%zmm2</code>
	<code>vpaddsw %zmm2,%zmm4,%zmm2</code>	<code>vpmulhw %zmm1,%zmm3,%zmm3</code>
	<code>vpcmpneqw %zmm3,%zmm2,%k1</code>	<code>vpsraw \$0xf,%zmm2,%zmm5</code>
	<code>kord %k1,%k0,%k0</code>	<code>vpcmpneqw %zmm3,%zmm5,%k1</code>
		<code>kord %k0,%k1,%k0</code>

Table 3.8: (TODO: write caption)

```
bool __builtin_add_overflow (type1 x, type2 y, type3 *sum);
```

```
bool __builtin_mul_overflow (type1 x, type2 y, type3 *prod);
```

These functions take three arguments: `x` and `y` are the two input operands, and `sum` or `prod` is a pointer to the variable that will hold the result of the addition or multiplication. The return value of these functions is a boolean that indicates whether an overflow occurred during the operation.

3.3.2.3 Evaluation

TODO: I have not done this benchmark before Compare how fast is the following:

1. overflow-checked, vectorized, int16
2. overflow-ignored, vectorized, int16
3. overflow-checked, scalar, int32
4. overflow-ignored, vectorized, int32
5. overflow-checked, scalar, int64
6. overflow-ignored, vectorized, int64

3.3.3 Overflow checking for floating points

To detect floating point overflow or imprecision, one approach is to enable floating point imprecision as a trap, then upon overflow, the interrupt `SIGFPE` is raised and the PC (program counter) will be redirected to its handler. The method can be programmed by using useful functions from the `fenv` library as follow [21]:

```
void signal_handler(int signal) {
    // handle fpe
}

void function() {
    std::signal(SIGFPE, signal_handler);
    std::feclearexcept (FE_ALL_EXCEPT);
    feenableexcept (FE_INEXACT | FE_INVALID);
    // do something
    fedisableexcept (FE_INEXACT | FE_INVALID);
}
```

In the `function()` block, first the `std::signal()` function is called to register the signal handler function for the `SIGFPE` interrupt. Next, the `std::feclearexcept()` function is called to clear any previously set exception flags in the status register. Then, `FE_INEXACT` and `FE_INVALID` is passed to the `feenableexcept()` function to enable inexactness and invalid exceptions. Once the floating-point exceptions are enabled, some computation can be performed. If some operation produces an inexact or invalid floating point number, `SIGFPE` is raised and a call to `signal_handler` is triggered. After the computation is completed, the `fedisableexcept()` function is called to disable the previously enabled exceptions.

But it is difficult to recover back from `SIGFPE`. By design, the propose usage of `SIGFPE` is to do some cleanup in the handler function, then exit the program gracefully. If the handler does not exit the program, after returning from the handler, the PC always points back to the instruction that caused `SIGFPE` and triggers `SIGFPE` again! However the goal is to discard the current progress and continue the program by using the `LargeInteger` algorithm. Although there could be potential workarounds, such as modifying the call stack and changing the return address, but implementing these solutions can be challenging, and introduces significant complexity to the codebase.

Alternatively we may read status registers and check if the imprecision bit is set:

```
bool function(matrix & tableau) {
    std::feclearexcept (FE_ALL_EXCEPT);
    if (fetestexcept(FE_INEXACT | FE_INVALID)) {
        // false for overflow, will be handle by its caller
        return false;
    } return true; // true for safe
}
```

Instead of registering floating point imprecision as interrupts, it clears the floating point status register and reads it afterwards using the `fetestexcept()` function, to check whether imprecision has ever occurred in previous computations. It then returns a boolean value to notify its caller whether or not the test for floating-point exceptions is positive, `true` for safe and `false` for overflown. In case of returning `false`, the caller will retry computation using `LargeInteger` accordingly.

3.3.3.1 Evaluation

In `x86_64` there are two status registers for floating points, the legacy `x87` status register for traditional scalar floating point operation, and the modern `mxcsr` register for SSE or AVX instructions. The source code of the library function `fetestexcept` indeed manipulates both registers [21]:

```
int fetestexcept(int excepts) {
    unsigned short status;
    unsigned int mxcsr;
```

```

    excepts &= FE_ALL_EXCEPT;

    /* Store the current x87 status register */
    __asm__ volatile ("fnstsw %0" : "=am" (status));

    /* Store the MXCSR register state */
    __asm__ volatile ("stmxcsr %0" : "=m" (mxcsr));

    return ((status | mxcsr) & excepts);
}

```

In the vectorized floating point scenario, the instruction for x87 status register is unnecessary and only the `mxcsr` register should be concerned. The hot loop is completely vectorized, and clang dispatches 128-bit SSE instructions for occasional scalar operations. This is because the SSE execution units can handle both single and double precision floating point arithmetic natively, whereas the legacy x87 floating-point instructions operates on an 80-bit internal format, requiring additional conversions and delay. Using SSE instructions reduces register pressure as well. The SSE units can leverage 16 XMM registers, but for x87 units there are only 8 floating point registers available. There could be even 32 XMM registers if such extension is implemented in the micro-architecture [22].

The benchmark, as illustrated in Figure ??, evaluates the performance of `fenv` functions alongside their respective revised versions. The modifications involves restricting operations solely to the manipulation of either the x87 status register or the `mxcsr` register. It indicates that it is significantly faster if we remove x87 status register related operations.

TODO: plot this

1. `fetestexcept` 8.15 ns
2. `fetestexcept_x87` 4.33 ns
3. `fetestexcept_mxcsr` 4.55 ns
4. `feclearexcept` 49.3 ns
5. `feclearexcept_local_x87` 45.5 ns
6. `feclearexcept_mxcsr` 9.42 ns

3.3.4 Comparing `int16_t` and `float`

Section 3.3.1, 3.3.2 and 3.3.3 have concluded that `int16_t` is superior than any other integer data types and `float` is better than `double`.

Benchmark (Figure 3.4) on the toy example reveals that `int16_t`'s spends a considerably higher percentage of runtime on overflow checking, compared with the floating point data types. This is consistent with the reasoning from previous chapters, where overhead for `float` is a one-time expense, but for `int16_t` it is always a portion of the total runtime.

It is worth noting that even though `int16_t` is faster than `float` in this benchmark,




Figure 3.4: TODO: replot and remove double

it is not sound to conclude that the `pivot` function implemented in `int16_t` will be faster than `float`. The toy example is different from the `pivot` function in many aspects, for example, number of memory load operations. In case of the actual `pivot` function, `float` may potentially outperform `int16`.

Chapter 4

Pivot

4.1 Implementation and Optimization

4.1.1 Matrix-wise transprecision

Transprecision techniques can be implemented at different levels of scale, such as element-wise, row-wise, and matrix-wise. The vectorized versions of `pivot` are implemented using the matrix-wise transprecision style. For `float`, the `mxcsr` register accumulates information regarding previous occurrences of overflow and imprecision, until it is cleared manually.

The element-wise method is not suitable in this scenario, as it defeats the purpose of vectorization. The row-wise method is not chosen either, due to that overflow is not likely to occur and the matrix is small. Reading the `mxcsr` can be considered as a expensive operation, comparing to the time spent on pivoting through the entire matrix. Avoiding wasted arithmetic instructions after overflow at the cost of more `mxcsr` reads is not a cost-effective trade-off.

Even though the `pivot` function using `int16_t` was implemented using the row-transprecision approach in previous works [6], it is modified to match with the matrix-wise transprecision style, in order to control differences between the `float` counterpart. After the overflow checking instructions, instead of a branch instruction pointing to the overflow handler, the boolean operator OR is applied between that overflow checking result register and a overflow flag.

4.1.2 Double buffering

Double buffering is a powerful technique to address the issue of data pollution caused by overflown data. When matrix-wise computing is carried out, a potential overflow can contaminate the input matrix and write meaningless results into it. It is difficult to recover from overflown results, making it impossible to dispatch the same input matrix to a algorithm of higher precision and defeating the purpose of transprecision computing.

Double buffering addresses this challenge by allocating two distinct pieces of memory, one for the input matrix and the other for output matrix. The input matrix is read-only, while the output matrix allows both reading and writing. This separation of data storage ensures that the input matrix remains unpolluted by overflowed data and that any potential overflow is encapsulated within the separate output matrix.

While making a copy of the input matrix is a simple and easy solution for protecting the original data from pollution, double buffering is a superior technique due that it does not introduce additional memory operations. With double buffering, every operand requires a load operation from the input matrix, and every result takes a store operation to be written into the output matrix. This is the minimal amount of memory operation required for arithmetics. Zen4 is capable of loading one ZMM vector per cycle and storing one ZMM vector per two cycles. In comparison, it can do two multiplication or addition of ZMM vectors every cycle. The discrepancy in throughput between computing and IO suggests that memory copy very expensive and inefficient.

4.1.3 Alignment and padding

4.1.4 Reduce number of matrix index computation

By placing the pivot row as the first row in the matrix,

In fact, the pivot function can be optimized to only compute index once (See Section 4.1.3) by placing the pivot row as first row.

4.1.5 Loop unrolling

4.1.6 Vector size specialization

Chapter 5

Conclusion and Future work

`half` is a 16-bit floating point data structure, consistent of 1-bit sign, 5-bit exponent and 10-bit mantissa (Figure 5.1). The `halfi` data type can be defined accordingly, and overflow would not occur for xx% of the cases.

Now `half` has already been supported widely on GPUs for AI applications. It provides performance benefits over higher precision floating point formats, for the exact same reason that extra precisions are a waste of resources, and more number can be computed once with a shorter data type. If 16-bit floating point is integrated into future AVX extensions, it would further improve the performance of `pivot`. It is expected to be slightly slower than the `int16_t` overflow-ignored version, and the runtime can be reduced to about 50% comparing to the `pivot` function implemented using `float` in this report.




Figure 5.1: IEEE 754 half precision floating point (16 bits). [13]

Bibliography

- [1] Google Benchmark. <https://github.com/google/benchmark>.
- [2] Mikolaj Bojanczyk and Joël Ouaknine. A simple and practical linear-time algorithm for presburger arithmetic. 2004.
- [3] LLVM Contributors. Mlir llvm. <https://mlir.llvm.org/docs/Dialects/Affine/>, 2023.
- [4] LLVM Contributors. Mlir llvm. <https://mlir.llvm.org/>, Accessed on 2023-04-03.
- [5] Grosser et al. Fast linear programming through transprecision computing on small and sparse data. In *Proceedings of the ACM on Programming Languages, Volume 4*, 2020.
- [6] Pitchanathan et al. Fpl: fast presburger arithmetic through transprecision. In *Proceedings of the ACM on Programming Languages, Volume 5*, 2021.
- [7] Tobias Gysi, Tobias Grosser, Laurin Brandner, and Torsten Hoefer. A fast analytical model of fully associative caches. In *Proceedings of the 40th ACM SIGPLAN Conference on Programming Language Design and Implementation*. ACM, jun 2019.
- [8] TODO. Todo, TODO.
- [9] TODO. Todo, TODO.
- [10] TODO. Todo, TODO.
- [11] TODO. Todo, TODO.
- [12] TODO. Todo, TODO.
- [13] TODO. Todo, TODO.
- [14] TODO. Todo, TODO.
- [15] TODO. Todo, TODO.
- [16] TODO. Todo, TODO.
- [17] TODO. Todo, TODO.
- [18] TODO. Todo, TODO.

- [19] TODO. Todo, TODO.
- [20] TODO. Todo, TODO.
- [21] TODO. Todo. TODO, TODO.
- [22] TODO. Todo, TODO.