Fast `pivot` Function for Presburger Library through Vectorization and Integer Arithmetic in FPU

Zhou Qi

4th Year Project Report
Computer Science
School of Informatics
University of Edinburgh

2023

Abstract

This report presents a fast implementation for the core function `pivot` of a math library in MLIR by performing vectorized integer arithmetics in FPU. The hot loop of the `pivot` function performs overflow-checked multiplication and addition on each element of an input matrix of low dimension and mostly small value items. MLIR's upstream uses element-wise transprecision, from `int64_t` to `LargeInteger`. This can be improved if hardware resources are utilized effectively, by taking advantage of SIMD, and reducing the bit width for every element. The current approach cannot be automatically vectorized by compiler, and `int64_t` has a much larger bit width than what is typically used for most of the items in the matrix. Additionally, extra arithmetics are required to perform overflow checking for integers, resulting in significant overhead. This report "innovates" the `int23_t` datatype, a 23-bit integer datatype created from the 23-bit mantissa of a 32-bit floating point, to address these issues. `int23_t` overflow can be captured as floating point imprecision by a status register, making overflow awareness almost free. On a selected 30-row by 20-column representative input matrix, the runtime of is reduced from 550 ns to 26 ns, achieving 20 times speedup.

Research Ethics Approval

This project was planned in accordance with the Informatics Research Ethics policy. It did not involve any aspects that required approval from the Informatics Research Ethics committee.

Declaration

I declare that this thesis was composed by myself, that the work contained herein is my own except where explicitly stated otherwise in the text, and that this work has not been submitted for any other degree or professional qualification except as specified.

(Zhou Qi)

Acknowledgements

Table of Contents

Chapter 1

Introduction

MLIR, Multi-Level Intermediate Representation, is a infrastructure for building reusable and extensible compilers. Its aim is to reduce fragmentation in domain specific languages and heterogeneous hardwares [6]. Its Presburger library provides polyhedral compilation techniques to make dependence analysis and loop optimization [5] and cache modeling [12]. Presburger arithmetics involves determining whether conjunction of linear arithmetic constraints is satisfiable [4], and can be solved using the simplex method of linear programming, with its core function `pivot` is the main performance bottleneck [10].

The `pivot` function involves two multiplication and one addition operation on every element in a matrix. Notably, the input matrices for this library tend to exhibit characteristics of small values and low dimensionality. For example, 90% of test cases work with 16-bit integers that never overflow, and 74% of isl's runtime is spent on test cases that we can compute using 16-bit integers and matrices with at most 32 columns [11]. These properties can be leveraged to take advantage from modern micro-architectural hardware resources, thereby accelerating the process.

Currently, the source code in MLIR upstream adopts a nested for-loop to iterate through every element of the matrix in a transprecision manner. Each number in the matrix can either be `int64_t` or `LargeInteger`. The algorithm starts by using `int64_t`, in case of overflow, it switches to the `LargeInteger` version. This approach is computationally expensive and inefficient, for the following reasons:

1. `int64_t` has a much larger bit width than what is typically used for most of the elements in the matrix,

2. the compiler is not capable of automatically generating vectorized instructions to further optimize the process,

3. overflow is checked manually through additional arithmetic operations.

To propose a faster alternative of the `pivot` function, we could consider constructing a new `pivot` algorithm that satisfies the following conditions:

1. Utilize SIMD: preliminary benchmark (Section 3.1.3) indicates at least 10 times performance improvement.

2. Use small bit width for every element: reducing bit width by half doubles the amount of numbers packed into a single vector register, and essentially reduces the instruction count by half (See Table 3.4).

3. Fast overflow checking: for integers, overflow has to be checked manually and this introduces at least 65% overhead towards total runtime (Section 3.3.2.1). This is because the x86 architecture does not provide status registers to indicate integer overflown. However, there is one for floating points, making floating points overflow detection almost free.

Previously there was an attempt to vectorize `pivot` that utilizes `int16_t` and targets matrices with 32 columns or less [11]. This approach offers the advantage of being able to pack a row of 32 elements into a single `AVX-512` register and addresses issues 1 and 2. However, overflow is still checked manually, causing 4 times more instruction count (Section 3.3.2.1). Additionally, this approach introduces a new disadvantage, `AVX-512` extensions are required for the support for vectorized `int16_t`, this is very rare among CPU manufactured in the last decade (Section 2.3).

An alternative approach is to do 23-bit or 52-bit integer operations using float (32-bit floating point) or double (64-bit floating point) respectively. Though floating points are notorious for precision issues, they are reliable when representing integers that fit inside their mantissa, 23 bits for float and 52 bits for double [1]. When the result of some integer computation exceeds the bit size of the mantissa, floating point imprecision almost always occurs and a status register will be set automatically (Section 3.3.3). Comparing to `int16_t`, even though vector size is sacrificed as there does not exist support for 16-bit floating point `half`, using floating points could still potentially be faster, because overflow checking overhead can be significantly reduced. With floating points, the cost of overflow checking is the time spent on resetting the status register once at the beginning of a sequence of calls to `pivot`, plus reading it in each `pivot` call. Even though reading the register takes 5 ns and resetting it costs 10.5 ns (Figure 3.5), the effective total overhead can be less than 1 ns. The average cost of resetting per `pivot` can be treated as negligible, while the superscalar and out-of-order execution pipeline hides the latency of reading the status register with pending memory operations. Moreover, floating points offer better compatibility with old computers than `int16_t`. Vector `float` or `double` code can be executed on CPUs with `AVX-2`, the predecessor of `AVX-512`. Almost every x86 CPU from the last decade supports this extension. It can also be adapted to fit even older `AVX-128` and `SSE` CPUs with minimal change in source code.

This report will first analyze the capability of modern CPU micro-architecture, especially `Zen4`, through a matrix element-wise fused-multiply-add toy example under the various configurations regarding vectorization methods, matrix data

---

[1]IEEE 754 specification is introduced in Section 2.4.1

structures, element data types and data widths (Section 3).

It is discovered that optimal performance can be achieved by selecting clang builtin vector type as vectorization methods and use flat list as matrix data structure. However, it is quite difficult to decide whether `int16_t` or `float` is better, because the former benefits form bigger vector size and less instruction count, while the latter has minimal overhead on overflow checking.

Then two detached versions of `pivot` function from the Presburger library are built from the most optimal configurations derived form the toy example, one using `int16_t` and the other using `float`. Some further optimizations were made by inspecting `perf` reports and assembly, including:

1. Implementing matrix-wise transprecision computing
2. Double buffering
3. Aligning the matrix to the size of vector registers
4. Reducing number of matrix index computation
5. Specialization for different row sizes

After applying these optimizations, a benchmark is setup targeting a selected 30-row matrix. The column of the matrix can be configured to be 8, 16, 24 or 32. In the case of 24 columns, which is considered to be the most representative case [10], it is discovered that:

1. `float` does not provide a performance advantage over `int16_t`. Specifically, `float` takes 26 ns, while `int16_t` completes 6 ns ahead. Nevertheless, both of them significantly outperform the upstream scalar implementation, which renders the 6 ns gap trivial.

2. Additionally, `float` offers substantial compatibility advantage over `int16_t` for vast amount of non-`AVX-512` platforms.

Chapter 2

Background

2.1   Linear Programming and Simplex Algorithm

Linear programming is a mathematical optimization technique used to model and find the best possible solution to a problem, given a set of constraints and an objective function to maximize or minimize. Its canonical form consists of a maximizing the objective function: $Z = c_1 x_1 + c_2 x_2 + ... + c_n x_n$, subjecting to the constraints:

$x_1 ... x_n \geq 0$

$a_{11} x_1 + a_{12} x_2 + ... + a_{1n} x_n = b_1$

$a_{21} x_1 + a_{22} x_2 + ... + a_{2n} x_n = b_2$

...

$a_{m1} x_1 + a_{m2} x_2 + ... + a_{mn} x_n = b_m$,

where $x_1 ... x_n$ are the variables, $c_1 ... c_n$ are the coefficients of the objective function, and non-negative $a_{11}, a_{12} ... a_{21} ... a_{m1} ... a_{mn}$ together with $b_1 ... b_m$ encodes the constraints of the problem in a matrix [10].

The simplex algorithm is an iterative approach to find $x_1 ... x_n$ that maximizes the objective function while satisfying constraints at the same time. The matrix goes through a sequence of transformations, until the solution appears or the solution is not feasible. The transformations are called "pivot", and it solves the linear equation at pivot row for the variable at pivot column:

$$
\begin{array}{lccl}
\text{Pivot item} & \alpha & \rightarrow & \frac{1}{\alpha} \\
\text{Rest of the pivot row} & \beta & \rightarrow & -\frac{\beta}{\alpha} \\
\text{Rest of the pivot column} & \gamma & \rightarrow & \frac{\gamma}{\alpha} \\
\text{Other entries} & \delta & \rightarrow & \delta - \frac{\beta\gamma}{\alpha}
\end{array}
$$

The pivot transformation involves division, and thereby produces rational numbers. Rationals of base 10 cannot be expressed as binary floating points precisely due to inaccuracy caused by potential rounding. In addition, division is an expensive operation comparing to addition or multiplication. This issue can be addressed through demoralizing each row by multiplying each row with their common

denominator [10]:

$$
\begin{array}{lccl}
\text{Pivot item} & \alpha & = & \frac{\alpha_n}{d_p} \\
\text{Rest of the pivot row} & \beta & = & \frac{\beta_n}{d_p} \\
\text{Rest of the pivot column} & \gamma & = & \frac{\gamma_n}{d_i} \\
\text{Other entries} & \delta & = & \frac{\delta_n}{d_i}
\end{array}
$$

where d_p is the denominator of the pivot row, d_i is the denominator of the i^{th} row. After demoralization α_n, β_n, γ_n, and δ_n, becomes α, β, γ, and δ.

Substituting into the processing formula, the pivot row becomes:

$$
\begin{array}{lccc}
\text{Pivot item} & \frac{\alpha_n}{d_p} & \rightarrow & \frac{d_p}{\alpha_n} \\
\text{Rest of the pivot row} & \frac{\beta_n}{d_p} & \rightarrow & -\frac{\beta_n}{\alpha_n}
\end{array}
$$

This transformation effectively becomes swapping α_n with d_p and negating every non-pivot-column item.

Likewise, the non-pivot rows are transformed as the following:

$$
\begin{array}{lccc}
\text{Rest of the pivot column} & \frac{\gamma_n}{d_i} & \rightarrow & \frac{\gamma_n d_p}{a_n d_i} \\
\text{Other entries} & \frac{\delta_n}{d_i} & \rightarrow & \frac{\delta_n a_n - \beta_n \gamma_n}{a_n d_i}
\end{array}
$$

and can be implemented in these procedures:

1. Update row denominator: $d_i' = d_i a_n$,
2. Multiply non-pivot-column items by a_n,
3. Subtract $\beta_n \gamma_n$ to every non-pivot row.

2.2 Presburger library

The Fast Presburger Library (FPL) paper collected 465,460 representative linear problems encountered during analyzing linear programs in cache analytical modeling, polyhedral loop optimization, and accelerator code generation. It is found that most of the constraint matrices are low in dimensionality and small in the value of each element [10]. Specifically, more than 99% of the coefficients require less than 10 bits and 95% of them are less than 20 columns (Figure 2.1). Thus, most of the rows fit inside a 512-bit vector register of 32 `int16_t` elements, and a row operation can be done in a single instruction.

However, in rare and corner cases, there can be larger coefficients up to 127 bits. Practically, the upper bound of coefficient size is unknown, making it required to have arbitrary precision arithmetic `LargeInteger` as a backup. Also, the maximum observed column count is 28 and there is not a certain maximum column count as well.

The FPL paper presents a 3-way transprecision implementation for the Presburger library's simplex solver using the algorithm described in Section 2.1. It starts from row-wise vectorized `int16_t`, and will switch to element-wise scalar `int64_t` or element-wise scalar `LargeInteger` in case of overflow, as illustrated in Figure 2.2. But unfortunately the MLIR upstream only presents a 2-layer transprecision, consisting of element-wise scalar operation using `int64_t` and `LargeInteger`. The `int16_t` version is not merged with the upstream for two reasons:

(a) Value Size - Bitwidth of coefficients in binary representation

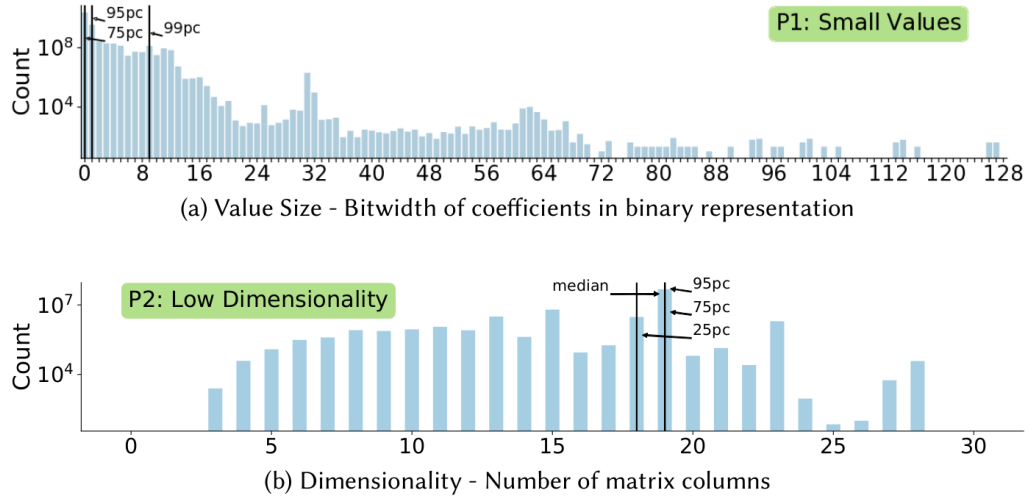(b) Dimensionality - Number of matrix columns

Figure 2.1: Linear programming problems for program analysis exhibits unique characteristics of small value size and low dimensionality [10].

1. `int16_t` vectors require `AVX-512` ISA extension, but hardware support is rare (Section 2.3).
2. Despite the `int16_t` version is fast, overflow checking overhead is 65% [11]. Using floating points could significantly reduce this overhead and potentially be faster (Section 2.4).

2.3 Modern CPU micro-architecture

A recent trend in x86-64 architecture's development is to include `AVX-512` instruction set architecture (ISA) extension. `AVX-512` succeeds `AVX-2`, the vector width is increased from `AVX-2`'s 256 bits to 512 bits. `AVX-512` also provides new instructions, for example, `int16_t` saturated addition.

2.3.1 Intel

Even though its specification was released by Intel in 2013, it had been unpopular [27], as it did not bring practical performance improvements. The primary reason was that it consumed a lot more power than usual, causing severe overheating. The micro-architecture Skylake from Intel, and its `AVX-512` enabled counterpart Skylake-X is a classic example. Skylake provides 2 256 bits FMA `AVX-2` execution units[1] and Intel provides 2 512-bit `AVX-512` FMA units by fusing the existing `AVX-2` units into a `AVX-512` unit, then introduces an additional FMA `AVX-512` unit [7]. The additional `AVX-512` unit increases the heat flux density of the chip, causing server thermal throttling issues.

Intel attempted to mitigate this problem by introducing the "`AVX-offset`" mode.

---

[1]Fused-multiply-add (FMA) execution units are a type of floating point execution units, capable of doing addition, multiplication or both in a single instruction. See Section 2.4.2.

Figure 2.2: The The architecture of FPL. The focus of this report is the "Simplex" method of the "Core Algorithms", since it is the main performance bottleneck. [11]

When a workload involving `AVX-512` instructions is encountered, the CPU automatically enters the `AVX-offset` mode and reduces its clock frequency [16]. This solution only works in theoretical benchmarks where `AVX-512` instructions presents in large bulk, but in practice it is more common to have a mix of control flow, scalar, `SSE` and `AVX-512` instructions. The clock frequency of executing those non-`AVX-512` instructions is decreased together with `AVX-512` instructions, causing many workloads could run faster with disabled `AVX-512` and higher clock frequency [15].

OptionalTODO: alderlake disabled avx-512 for big.LITTLE

2.3.2 AMD

AMD recently decides to add support for `AVX-512` in their latest micro-architecture Zen4. It has slightly less computing power comparing with Intel, but much more efficient. Zen4 can be considered as modernized version of Zen3 or Zen2, where Zen2 and Zen3 supports `AVX-2` by providing 2 FADD units [2] and 2 FMA units of 256-bit width [3]. Zen4 "double-pumps" these existing circuits to create a single 512-bit FADD and a single 512-bit FMA, without introducing any new arithmetic units [15]. Zen2 and Zen3 are reputable for its high performance per watt [18], and Zen4 would be better with its more advanced lithography [15].

Additionally, rebuilding existing software to target `AVX-512` may bring slight performance improvement. One benefit of `AVX-512` is that it reduces front-end pressure. In the case of the Zen4 micro-architecture, though the back-end is possible to commit 2 `AVX-2` FADD and 2 `AVX-2` FMA every cycle, the front-end

---

[2]Floating-point add units (FADD) can execute addition instructions only. They may be considered as simplified FMA

Figure 2.3: IEEE 754 specification for single (32 bits) and double (64 bits) precision floating point [17]. In some literature "mantissa" is referred as "fraction".

has to dispatch 4 instructions per cycle, which is quite difficult. The equivalency in `AVX-512` only takes 2 instructions, this is much more likely to be sustained by the frontend [15].

2.4 Floating points

2.4.1 IEEE 754

IEEE 754 is the standard for representing and manipulating floating-point numbers in modern x86 computers. The standard defines several different formats for representing floating point numbers, the most common ones are 32-bit single precision (float) and 64-bit double precision (double). For each format, it specifies how many bits are used to represent the sign, exponent, and mantissa.

For float and double, the sign bit is a single bit that indicates whether the number is positive or negative. As Figure 2.3 shows, there are 8 bits and 11 bits for exponent in float and double respectively, to represent represents the order of magnitude. The remaining 23 bits in float and 52 bits in double are mantissae, to store the fractional part of the number. The value of a floating point number can be computed through this formula: $(-1)^s * 2^{(e-B)} * (1 + f)$ where s is sign, e is exponent, f is mantissa and B is a constant bias value: 127 for float, 1023 for double.

Figure 2.3 provides an example of `float` by presenting `0.15625` in binary form:
```
sign     = 0b0          -> 0
exponent = 0b01111100 -> 0b01111100 - 127 = 124 - 127 = -3
mantissa = 0b01         -> 0b1.01 = 1.25
```
Substituting the sign, exponent and mantissa into the formula, we get:
```
-1^0 * 2^(-3) * 1.25 = 0.15625.
```

2.4.2 Fused-multiply-add

After doing floating point arithmetic, it is required to normalize the result of floating-point arithmetic before it can be used further. However, by feeding the result of a floating-point multiplication (FMUL) directly into the floating-point addition (FADD) logic without the need for normalization and rounding in between, a fused multiply-add (FMA) operation is effectively created: $Y = (A * B) + C$, where A, B and C are the operands, Y is the result [19].

FMA saves cycles and reduces accumulation of rounding errors, while at the same time not adding significant complexity to the circuit. A FMA execution unit is capable to do FMUL, FADD, and FSUB as well:

Addition: $Y = (A * 1.0) + C$
Multiplication: $Y = (A * B) + 0.0$
Subtraction: $Y = (A * -1.0) + C$

This is a useful feature in many numerical computations that involve simultaneous multiplication and addition operations, such as dot produpreliminaryct and matrix multiplication. Since the `pivot` function performs multiplication and addition between the pivot row, some constant value and each row in the matrix, the performance of FMA is critical to the overall efficiency of the algorithm.

2.4.3 Representing "`int23_t`" and "`int52_t`" using floating points

There is a common stereotype that floating-point numbers are unreliable and likely to be imprecise, and are often illustrated in popular memes (as shown in Figure 2.4). However, when storing integer values inside floating points, floating points can be quite reliable.

Specifically, given that the mantissa part of a float consists of 23 bits, inexactness never occur when representing integers less than 23 bits (ignoring the sign bit). Furthermore, in case of of an integer overflow, floating point imprecision almost always occurs and a corresponding status registers is set. The same concept applies to double data types, which have a mantissa consisting of 52 bits.

This mechanism is reliable, as floating-point inexactness always implies integer-inexactness. For a integer value with bit width greater than the mantissa size, floating point rounding is triggered in order to fit most significant bits of the integer in the mantissa, and then adjust the order of magnitude in the exponent accordingly. The lower bits of the mantissa are truncated, therefore causing imprecision.

In some rare cases, integers longer than the size of mantissa can be represented in floating points precisely. An examples of such numbers are large powers of 2, like
`2^30 = 0x40000000`. Its binary representation in `float` is:

```
Sign     = 0b0         -> 0
Exponent = 0b10011011 -> 0b10011011 - 127 = 155 - 127 = 28
Mantissa = 0b0         -> 0b1.0 = 1
```

Figure 2.4: A floating point meme: `0.1 + 0.2 = 0.30000000000000004`

Despite being greater than the size of the mantissa, they are normalized rather then being rounded, and therefore does not break the mechanism of representing integer in floating-points.

Throughout history, floating point processing units in GPUs have been utilized for fast integer arithmetic, for example, modular exponentiation [20] and RSA algorithms [21], because often the architecture of GPU prioritizes floating points performance rather than integers.

2.5 Google Benchmark

Google benchmark is a library to measure the performance of a code snippet. It provides unit-test like interfaces to setup benchmarks around a code snippet [8]. The given example from https://github.com/google/benchmark is self-explanatory for its usage:

```
#include <benchmark/benchmark.h>

static void BM_SomeFunction(benchmark::State& state) {
    // Perform setup here
    for (auto _ : state) {
    // This code gets timed
    SomeFunction();
    }
}
// Register the function as a benchmark
BENCHMARK(BM_SomeFunction);
// Run the benchmark
BENCHMARK_MAIN();
```

The library first starts a timer, repeatedly executes its core loop: `for (auto _ :`

`state)` `...` multiple times then pauses the timer. This method ensures that the results are consistent and minimizes the overhead required for recording the timing information.

Executing the benchmarks will not only report both elapsed real time and CPU time, but also much other useful information to help reduce variance.

```
Running ./build/example
***WARNING*** CPU scaling is enabled, the benchmark real time
measurements may be noisy and will incur extra overhead.
Run on (32 X 5800.00 MHz CPU s)
CPU Caches:
    L1 Data 32 KiB (x16)
    L1 Instruction 32 KiB (x16)
    L2 Unified 1024 KiB (x16)
    L3 Unified 32768 KiB (x2)
Load Average: 8.10, 5.14, 1.14
------------------------------------------------------------
Benchmark                    Time            CPU   Iterations
------------------------------------------------------------
BM_SomeFunction           18.5 ns        18.5 ns     37935734
```

The warning: "CPU scaling is enabled, the benchmark real-time measurements may be noisy and will incur extra overhead." is saying that CPU clock frequency is not consistent. It can be dynamically determined by the governor algorithm, according on the system's needs. For example, with the `performance` governor, the OS locks the CPU to the highest possible clock frequency, specified at `/sys/devices/system/cpu/cpu*/cpufreq/scaling_max_freq`, while the `ondemand` governor will push the CPU to the highest frequency on demand and then gradually reduce the frequency as the idle time increases [13].

However, it is also dependent on the manufacture and other hardware constraints. By default, both Intel (Turbo Boost) and AMD (Precision Boost Overdrive) have support for raising clock frequency, beyond the control of the governor [22]. On the other hand, CPUs have self-protecting thermal throttling mechanisms that reduces its clock frequency and voltage when it is too hot.

The benchmark mentioned in this report were performed on a AMD 7950x desktop computer. The computer system went through the following these setups for consistent results:

1. Set the governor to `performance`,
2. Disable AMD Precision Boost Overdrive (or Intel Turbo Boost),
3. Lock clock frequency at a 4.5 GHz, or any desired and feasible value,
4. Make sure heat dissipation is working properly.

2.6 `llvm-mca`

`llvm-mca`, LLVM Machine Code Analyzer, is a tool to analyze performance of executing some instructions on a specific CPU micro-architecture, according to scheduling information provided by LLVM [9].

By supplying `llvm-mca` with a piece of assembly code and the target micro-architecture codename, `llvm-mca` reports various metrics to indicate how fast the given instructions will execute on the specified micro-architecture. It first summarizes the instruction per clock (IPC) and throughput of the entire instruction block, then gives detailed information about each instruction, including number of uOps, latency, throughput, potential load, store and side effects. `llvm-mca` also reports resource pressure in terms of arithmetic units and memory load or store units. When the optional `-timeline` flag is prompted, `llvm-mca` illustrates a timeline view of the analyzed code, showing how instructions progress through the pipeline stages of the target processor. The timeline helps understand the capability of complicated out-of-order superscalar architecture.

An example is provided below. The analysis from `llvm-mca` indicates that a `znver2` (Zen2) CPU can repeatedly execute a combination of `vmovaps`[3] and `vfmadd213ps`[4] instructions every 1.3 cycle. This translates to 2.74 instructions per cycle. The output from `llvm-mca` is slightly modified and truncated to fit into limited page size.

```
$ llvm-mca-15 -timeline -mcpu=znver2 ./x.s
Iterations:        100
Instructions:      400
Total Cycles:      146
Total uOps:        400

Dispatch Width:    4
uOps Per Cycle:    2.74
IPC:               2.74
Block RThroughput: 1.3

Instruction Info:
[1]: #uOps
[2]: Latency
[3]: RThroughput
[4]: MayLoad
[5]: MayStore
[6]: HasSideEffects (U)

[1] [2] [3]  [4] [5] [6] Instructions:
 1   8  0.33  *              vmovaps  (%rdx,%rsi,4), %ymm0
```

---

[3]`vmovaps` can be either vector float load or store instruction, depending on how operands are structured.

[4]`vfmadd213ps` is the instruction for fused-multiply-add.

```
1   8  0.33   *            vmovaps   (%rcx,%rax,4), %ymm1
1   12 0.50   *            vfmadd213ps  (%r8,%rdi,4), %ymm0, %ymm1
1   1  0.33        *       vmovaps   %ymm1, (%r9,%rax,4)


Resources:
[0]    - Zn2AGU0
[1]    - Zn2AGU1
[2]    - Zn2AGU2
...
[8]    - Zn2FPU0
[9]    - Zn2FPU1
[10]   - Zn2FPU2
[11]   - Zn2FPU3
[12]   - Zn2Multiplier


Resource pressure per iteration:
[0]   [1]   [2]   [3] [4] [5] [6] [7] [8]   [9] [10] [11] [12]
1.33 1.33 1.34  -   -   -   -   -  0.50  -   -   0.50  -


Resource pressure by instruction:
[0]   [1]   [2]   ... [8]   [11]  Instructions:
0.01 0.38 0.61 ...  -     -     vmovaps   (%rdx,%rsi,4), %ymm0
0.23 0.68 0.09 ...  -     -     vmovaps   (%rcx,%rax,4), %ymm1
0.27 0.24 0.49 ... 0.50  0.50  vfmadd213ps  (%r8,%rdi,4), %ymm0, %ymm1
0.82 0.03 0.15 ...  -     -     vmovaps   %ymm1, (%r9,%rax,4)


Timeline view:
                0123456789
Index 0123456789          012345


[0,0] DeeeeeeeeER    .     .    . vmovaps   (%rdx,%rsi,4), %ymm0
[0,1] DeeeeeeeeER    .     .    . vmovaps   (%rcx,%rax,4), %ymm1
[0,2] D=eeeeeeeeeeeER    .      . vfmadd213ps  (%r8,%rdi,4), %ymm0, %ymm1
[0,3] D============eER   .      . vmovaps   %ymm1, (%r9,%rax,4)
[1,0] .DeeeeeeeeE-----R   .     . vmovaps   (%rdx,%rsi,4), %ymm0
[1,1] .DeeeeeeeeE-----R   .     . vmovaps   (%rcx,%rax,4), %ymm1
[1,2] .D=eeeeeeeeeeeER    .     . vfmadd213ps  (%r8,%rdi,4), %ymm0, %ymm1
[1,3] .D============eER   .     . vmovaps   %ymm1, (%r9,%rax,4)
[2,0] . DeeeeeeeeE-----R   .    . vmovaps   (%rdx,%rsi,4), %ymm0
[2,1] . DeeeeeeeeE-----R   .    . vmovaps   (%rcx,%rax,4), %ymm1
...
```

However, when evaluating the identical assembly code on a more advanced micro-architecture `znver3` (Zen3), `llvm-mca` reveals reduction in IPC and throughput. This appears to be contradictory with both theoretical expectations and actual benchmarks. After submitting an issue [2], llvm maintainers explained that llvm's

scheduling information are hand-crafted using `llvm-exegesis`, a micro-benchmark tool. The issue was subsequently resolved after rerunning `llvm-exegesis` and confirming that `znver3` indeed has higher throughput as expected.

TODO: should I write a paragraph to emphacise that `llvm-mca` is unreliable?

Chapter 3

Experiments with Toy Example

The `pivot` function does multiply and add for each row in the matrix, therefore the performance of FMA a simple vector toy can be an effective indicator. This chapter reports performance analysis on simple toy examples that do vector add or vector FMA with various setups, including:

1. Vectorization method
   (a) Clang's automatic vectorization from scalar source code
   (b) Clang builtin vector datatype with occasional AVX intrinsics

2. Matrix data structure
   (a) Nested list
   (b) Flat list

3. Element data width
   (a) 16 bits: `int16_t`
   (b) 32 bits: `int32_t`, `float`
   (c) 64 bits: `int64_t`, `double`

4. Element data type
   (a) Integer
   (b) Floating point

3.1 Vectorization method

3.1.1 Clang's automatic vectorization

Clang is capable of generating vectorized instructions from scalar source code, using the flags `-O3 -march=native` on a platform with vector ISA enabled. Starting with an example (Listing 3.1), the simple `vec_add` function adds every element from two arrays and saves it to the third.

Source code

```
#define size 128
void vec_add(float* src1_ptr, float* src2_ptr, float* dst_ptr) {
    for (uint32_t i = 0; i < size; i += 1 ){
        dst_ptr[i] = src1_ptr[i] + src2_ptr[i];
    }
}
```

Assembly snippet of the hot loop, vectorization on

```
1458: c4 c1 7c 58 84 87 20    vaddps -0x1e0(%r15,%rax,4),%zmm0,%zmm0
145f: fe ff ff
1462: c4 c1 74 58 8c 87 40    vaddps -0x1a0(%r15,%rax,4),%zmm1,%zmm1
1469: fe ff ff
```

Assembly snippet of the hot loop, vectorization off

```
120d: d8 44 82 04    fadds  0x4(%rdx,%rax,4)
1211: d9 5c 81 04    fstps  0x4(%rcx,%rax,4)
1215: d9 44 86 08    flds   0x8(%rsi,%rax,4)
```

Listing 3.1: The vectorized and scalar binary are derived by compiling with flags -O3 -march=native and -O3 -march=native -mno-avx -mno-sse respectively, on a Zen4 computer with clang-15.

After compiling on a AVX-512 enabled computer and disassembling the binary, it is observed that clang automatically packs 16 float (512 bits) as a operand of the vaddps instruction.

Alternatively, vectorization could be disabled by adding the -mno-avx -mno-sse flags on top of -O3 -march=native. These two sets of flags guarantee that the binary are going to be equally optimized, with the only difference been whether vector instructions are generated or not. In this case, scalar instructions fadds, fstps and flds are selected.

3.1.2 Clang's vector datatype and AVX intrinsics

Another approach is to write source code with vectorization in mind in the first place. Clang provides extension that allows programmers to declare a new type that represents a vector of elements of the same data type. The syntax is typedef *ty* vec_ty __attribute__((ext_vector_type(*vec_width*))), where **vec_ty** is the name of vector type being defined, **vec_width** is its size and *ty* is the type of the elements in the vector. For example, typedef int16_t int16x32 __attribute__((ext_vector_type(32))) defines a 512-bit vector type of int16x32, consisting of 32 int16_t and fits inside an AVX-512 ZMM register.

After defining a vector datatype, a vector variable can be created by casting from a pointer of the target array. Then arithmetic operators can be applied between the vectors to performed element-wise operations. The previous `vec_add` example can be rewritten as the code snippet shown in Listing 3.2:

Source code

```
#define size 128
#define FloatZmmSize 16
typedef float floatZmm __attribute__((ext_vector_type(FloatZmmSize)));
void vec_add(float* src1_ptr, float* src2_ptr, float* dst_ptr) {
    for (uint32_t i = 0; i < size; i += FloatZmmSize){
        floatZmm src1Vec = *(floatZmm *)(src1_ptr + i);
        floatZmm src2Vec = *(floatZmm *)(src2_ptr + i);
        *(floatZmm *)(dst_ptr + i) = src1Vec + src2Vec;
    }
}
```

Listing 3.2: Comparing to the C++ source code in Listing 3.1, it is slightly more complicated to code with vector types.

3.1.3   Evaluation

When comparing the performance of code written with and without the vector type and examining their assembly, it has been discovered that the the automatic vectorization feature in clang can be unpredictable and may lead to undesired behaviors. It operates as a black box and may take a lot of effort to understand its mechanisms. One of the issues is that clang may select a suboptimal vector width.

Consider the `vec_fma` function in Listing 3.3, a slightly more complicated version of the previous `vec_add` example, where there are 3 input matrices and the element-wise operation is changed from addition to FMA. The disassembly reveals that clang decides to use FMA vector instructions of 128-bit width, but when vector size is constrained to a bigger width by defining a vector type, more optimal binary can be generated. Benchmark (Figure 3.1) shows that the vector type version is 6 times and 11 times faster than the automatic vectorization and vectorization disabled version respectively.

Scalar source code to be automatically vectorized by clang

```
void vec_fma(matrix<float> & mat_src1, matrix<float> & mat_src2,
             matrix<float> & mat_src3, matrix<float> & mat_dst) {
    for (int i = 0; i < row; i += 1) {
        for (int j = 0; j < col; j += 1) {
            float src1 = mat_src1.get(i,j);
            float src2 = mat_src2.get(i,j);
            float src3 = mat_src3.get(i,j);
            mat_dst.set(i,j,  (src1 * src2) + src3);
        }
    }
}
```

Assembly snippet of the hot loop

```
80e4: c5 fa 10 04 b2     vmovss (%rdx,%rsi,4),%xmm0
80f1: c5 fa 10 0c 81     vmovss (%rcx,%rax,4),%xmm1
80fe: c4 c2 79 a9 0c b8  vfmadd213ss (%r8,%rdi,4),%xmm0,%xmm1
810c: c4 c1 7a 11 0c 81  vmovss %xmm1,(%r9,%rax,4)
```

Source code written with clang's vector type

```
#define FloatZmmSize 16
typedef float floatZmm __attribute__((ext_vector_type(FloatZmmSize)));
void vec_fma(matrix<float> & mat_src1, matrix<float> & mat_src2,
             matrix<float> & mat_src3, matrix<float> & mat_dst) {
    float * src1_ptr = (float *) mat_src1.getItemPointer(0,0);
    float * src2_ptr = (float *) mat_src2.getItemPointer(0,0);
    float * src3_ptr = (float *) mat_src3.getItemPointer(0,0);
    float * dst_ptr  = (float *) mat_dst.getItemPointer(0,0);
    for (uint32_t i = 0; i < row * col; i += FloatZmmSize){
        floatZmm src1Vec = *(floatZmm *)(src1_ptr + i);
        floatZmm src2Vec = *(floatZmm *)(src2_ptr + i);
        floatZmm src3Vec = *(floatZmm *)(src3_ptr + i);
        *(floatZmm *)(dst_ptr + i) = src1Vec * src2Vec + src3Vec;
    }
}
```

Assembly snippet of the hot loop

```
82c0: 62 b1 7c 48 10 04 80    vmovups (%rax,%r8,4),%zmm0
82c7: 62 b1 7c 48 10 0c 81    vmovups (%rcx,%r8,4),%zmm1
82ce: 62 b2 7d 48 a8 0c 86    vfmadd213ps (%rsi,%r8,4),%zmm0,%zmm1
82d5: 62 b1 7c 48 11 0c 87    vmovups %zmm1,(%rdi,%r8,4)
```

Listing 3.3: The toy example performs vectorized fused-multiply-add operation on every item from 3 input matrices and save the result to the output matrix. The internal data structure of `matrix` is `std::vector`, and its member function `getItemPointer(r,c)` returns the pointer to the element at row r and column c. The key distinction between two vectorization approaches are the register types. XMMs are 128-bit registers, while ZMMs are 512-bit long.

In some cases clang could be even worse, it may fail to recognize vectorization patterns from element-wise loop operations, leading to more reduction in performance. In the `vec_fma` example, by changing the type signature from `float` to `int`, clang decides to dispatch scalar instructions for addition (`add`) and multiplication (`imul`) completely (Listing 3.4). Their vectorized equivalency `vpaddd` and `vpmulld` is 15 times more performant (Figure 3.1).


Figure 3.1: A benchmark for the element-wise multiply-add toy example on a 16 by 16 matrix with different vectorization techniques. The toy example written with vector types are 5 to 10 times faster than their automatically vectorized or scalar counterpart.

Scalar source code to be automatically vectorized by clang

```
void vec_fma(matrix<int> & mat_src1, matrix<int> & mat_src2,
             matrix<int> & mat_src3, matrix<int> & mat_dst) {
    for (int i = 0; i < row; i += 1) {
        for (int j = 0; j < col; j += 1) {
            int src1 = mat_src1.get(i,j);
            int src2 = mat_src2.get(i,j);
            int src3 = mat_src3.get(i,j);
            mat_dst.set(i,j,  (src1 * src2) + src3);
        }
    }
}
```

Assembly snippet of the hot loop

```
88b0: 44 8b 69 1c   mov     0x1c(%rcx),%r13d
88b4: 44 8b 62 1c   mov     0x1c(%rdx),%r12d
88b8: 45 0f af ee   imul    %r14d,%r13d
88bc: 45 0f af e6   imul    %r14d,%r12d
88c0: 45 01 fd      add     %r15d,%r13d
88c3: 45 01 fc      add     %r15d,%r12d
```

Vectorized source code using clang's vector type extension

```
#define IntZmmSize 16
typedef int intZmm __attribute__((ext_vector_type(IntZmmSize)));
void vec_fma(matrix<int> & mat_src1, matrix<int> & mat_src2,
             matrix<int> & mat_src3, matrix<int> & mat_dst) {
    int * src1_ptr = (int *) mat_src1.getItemPointer(0,0);
    int * src2_ptr = (int *) mat_src2.getItemPointer(0,0);
    int * src3_ptr = (int *) mat_src3.getItemPointer(0,0);
    int * dst_ptr  = (int *) mat_dst.getItemPointer(0,0);
    for (uint32_t i = 0; i < row * col; i += IntZmmSize){
        intZmm src1Vec = *(intZmm *)(src1_ptr + i);
        intZmm src2Vec = *(intZmm *)(src2_ptr + i);
        intZmm src3Vec = *(intZmm *)(src3_ptr + i);
        *(intZmm *)(dst_ptr + i) = src1Vec * src2Vec + src3Vec;
    }
}
```

Assembly snippet of the hot loop

```
8880: 62 b1 fe 48 6f 04 81   vmovdqu64 (%rcx,%r8,4),%zmm0
8887: 62 b2 7d 48 40 04 80   vpmulld (%rax,%r8,4),%zmm0,%zmm0
888e: 62 b1 7d 48 fe 04 86   vpaddd (%rsi,%r8,4),%zmm0,%zmm0
8895: 62 b1 fe 48 7f 04 87   vmovdqu64 %zmm0,(%rdi,%r8,4)
```

Listing 3.4: Comparing to source code from Listing 3.3, the only change is replacing float with int32_t. Clang fails to vectorize the scalar version, but writing vector type guarantees vectorization.

3.2 Matrix data structure

The most intuitive data structure of a matrix is a list of lists, where each list represents a row and a list of rows is a matrix. In C++ this can be represented using `std::vector<std::vector<T>>`, where T could be `float`, `double`, `int32_t`, etc. The std::vector class provides an intuitive interface for accessing and modifying elements, making it easy to write code with.

One potential drawback of nested `std::vector` is that it requires two indexing operations to access an element. An alternative implementation is to "flatten" a matrix into a single `std::vector`, by simply concatenating one row after another. To access a specific element, an index can be computed manually using the given row and column: `column_count * row + column`. This reduces half of the memory indexing operation at the cost of additional arithmetic. The differences between the two patterns are illustrated by a example provided in Table 3.5.

	Nested	Flat
Type	`std::vector<` `std::vector<int32_t>>`	`std::vector<int32_t>`

Structure in Memory	```vector of 4 = {```
	    `vector of 4 = {0, 0, 0, 0},`
	    `vector of 4 = {0, 0, 0, 0},`
	    `vector of 4 = {0, 0, 0, 1},`
	    `vector of 4 = {0, 0, 0, 0}`
	`}`
		`vector of 16 = {`
		    `0, 0, 0, 0,`
		    `0, 0, 0, 0,`
		    `0, 0, 0, 1,`
		    `0, 0, 0, 0`
		`}`

Accessing row 2, column 3	Index row first:
	`std::vector<int32_t>[2]`
	then index column:
	`std::vector<int32_t>[2][3]`
		Compute `i =`
		`col_count * row +`
		`col`
		`= 4 * 2 + 3 = 11,`
		then index once:
		`std::vector<int32_t>[11]`

Table 3.5: This is an example of a 4 by 4 matrix structured using nested list and flat list, to highlight their differences.

Empirically indexing costs more time than integer multiplication and addition, thereby improving performance. Both the toy example and the `pivot` function perform sequential load-compute-store operations on each row and each column, allowing the index of the next element to be computed by simply adding the step size or column size, further reducing memory overhead. In fact, the `pivot` function can be optimized to only compute index once (See Section 4.4) by placing the pivot row as first row. Benchmark (Figure 3.2) on the toy example confirms that when there are 16 rows, the nested vector matrix is about 8 ns faster than the flat matrix.

Figure 3.2: On a 16-row toy example doing `float` fused-multiply-add, selecting a flat list as the data structure of matrix is consistently faster than the nested list implementation.

3.3 Matrix element data type

3.3.1 Width

Since the numbers stored in the matrix are almost always less than 10 bits, using shorter data types can be more advantageous than longer ones because they allow more numbers to be packed into a single vector register (Figure 3.4). The number of instructions can be cut by half when data width is reduced to half, and less instruction count always leads to less execution time. Given that the Zen4 micro-architecture provides approximately same amount of execution units for both integers and floating points, it is reasonable to estimate that the execution time is inversely proportional to the bit width of data type. As confirmed in Figure 3.3, when overflow is ignored, `int32_t` and `float` costs nearly same amount of time, while `int32_t` and `double` costs double the amount of time than `int16_t` and `float` respectively.

Figure 3.3: A benchmark for toy example on a 16 by 16 matrix of `int16_t`, `int32_t`, `int64_t`, `float`, and `double`, with overflow checking turned on or off. The runtime information for overflow checked `int32_t` and `int64_t` is not available, because it is difficult to implement vectorized overflow checker for them. It is discovered that (1) runtime is reduced by 50% as the bit width of data type is cut by half, (2) overflow checking for integer is much more expensive then floating points.

	float	double	int16_t	int32_t	int64_t
512-bit units	1 512-bit FADD + 1 512-bit FMA		2 512-bit ALU		
256-bit units	2 256-bit FADD + 2 256-bit FMA		4 256-bit ALU		
Fused-multiply-add	Yes		No		only on lower 52 bits, no overflow exception
Saturated add	N/A		Yes	No	No
Multiply higher bits	N/A		Yes	No	No
SIMD Floating-Point Exceptions	Overflow, Underflow, Invalid, Precision, Denormal		No		
512-bit vector size	16	8	32	16	8
256-bit vector size	8	4	16	8	4
Overflow checking and time cost	single time overhead: read status register 4.5 ns clear status register 9.4 ns		additional arithmetic, about 4x more instructions	Must fallback to scalar code due to lack of saturated add and multiply higher bits.	

Figure 3.4: A summary of features and resources provided by the Zen4 micro-architecture for different data types [15] [3].

3.3.2 Overflow checking for integers

The x86-64 micro-architecture provides the `seto` instruction to set some byte to 1 if overflow occurred as a result of integer arithmetic. However `seto` only

works for scalar operations, there does not exist instruction or status register to indicate whether a previous vector add or multiply instruction produces overflown results or not. Therefore, overflow has to be checked manually by some additional vector instructions, this would slow down the computation to some extent. Or alternatively or arithmetics have to be carried out on each element individually in a scalar manner,resulting in even worse performance.

One advantage of `int16_t` is that it can be used with `AVX-512`'s saturated add and multiply higher bits vector instructions (Figure 3.4), making it possible and convenient to write vectorized and overflow-aware code. However, Zen4's implementation of `AVX-512` extension does not provide equivalent instruction for `int32_t` or `int64_t`, and therefore must be processed as scalar values.

3.3.2.1 Implementation of vectorized `int16_t` overflow checking

By comparing the result of a conventional addition and saturated addition, it indicates whether an addition has gone overflown or not. In case of overflow, with saturated add, the result always retains at the maximum possible value of `int16_t`: `0x7FFF`, while the result of a conventional add is always smaller, because the overflowed output from conventional addition can't go all the way around and become `INT16 MAX` again. In the two's complement binary form for integer, the overflow sum is "trapped" in the negative number space. For example:

```
INT16_MAX + 1 = INT16_MIN = -32768,
INT16_MAX + 2 = -32767,
...
INT16_MAX + INT16_MAX = -2,
```

For multiplication, as two 16-bit numbers produces 32-bit products but only lower 16 bits can be stored, overflow can be detected by checking whether any of the upper 16 bits are set.

Inspecting these approaches from a instruction-level perspective (Table 3.6), when overflow is ignored, both add and multiply takes 1 instruction, `vpaddw`[1] and `vpmullw`[2]. To obtain and process overflow-related information, an additional computation instruction `vpaddsw`[3] or `vpmulhw`[4] is required, followed with 2 or 3 comparison, shuffling and branch instructions: `vpsraw`[5], `vpcmpneqw`[6] and `kord`[7]. By enabling overflow checking, it brings 4 to 5 times more instruction count and 65% more runtime [11].

---

[1] Vector add for `int16_t`
[2] Vector multiply lower half bits for `int16_t`
[3] Vector saturated add for `int16_t`
[4] Vector multiple higher half bits for `int16_t`
[5] Shift packed data right arithmetic
[6] Compare packed data for equal
[7] Bitwise logical OR masks

	Addition	Multiplication
Overflow ignored	`vpaddw %zmm4,%zmm2,%zmm3`	`vpmullw %zmm1,%zmm3,%zmm2`

Overflow aware	`vpaddw %zmm4,%zmm2,%zmm3`
	`vpaddsw %zmm2,%zmm4,%zmm2`
	`vpcmpneqw %zmm3,%zmm2,%k1`
	`kord %k1,%k0,%k0`
		`vpmullw %zmm1,%zmm3,%zmm2`
		`vpmulhw %zmm1,%zmm3,%zmm3`
		`vpsraw $0xf,%zmm2,%zmm5`
		`vpcmpneqw %zmm3,%zmm5,%k1`
		`kord %k0,%k1,%k0`

Table 3.6: This table highlights the difference in terms of instruction count when overflow checking is enable or disabled for vectorized `int16_t`. Pink instructions are essential components as they compute the anticipated arithmetic results, while overflow information are provided by yellow and cyan instructions.

3.3.2.2   Implementation of scalar `int32_t` and `int64_t` overflow checking

TODO:Not sure if this section is necessary or not?

Clang's language extension provides functions to perform overflow-checked integer arithmetics:
`bool __builtin_add_overflow (type1 x, type2 y, type3 *sum);`
`bool __builtin_mul_overflow (type1 x, type2 y, type3 *prod);`
These functions take three arguments: `x` and `y` are the two input operands, and `sum` or `prod` is a pointer to the variable that will hold the result of the addition or multiplication. The return value of these functions is a boolean that indicates whether an overflow occurred during the operation.

3.3.3   Overflow checking for floating points

To detect floating point overflow or imprecision, one approach is to enable floating point imprecision as a trap, then upon overflow, the interrupt SIGFPE is raised and the PC (program counter) will be redirected to its handler. The method can be programmed by using useful functions from the **fenv** library as follow [23]:

```
void signal_handler(int signal) {
    // handle fpe
}


void function() {
    std::signal(SIGFPE, signal_handler);
    std::feclearexcept (FE_ALL_EXCEPT);
    feenableexcept (FE_INEXACT | FE_INVALID);
    // do something
    fedisableexcept (FE_INEXACT | FE_INVALID);
}
```

In the `function()` block, first the `std::signal()` function is called to register the signal handler function for the SIGFPE interrupt. Next, the `feclearexcept`

function is called to clear any previously set exception flags in the status register. Then, FE_INEXACT and FE_INVALID is passed to the feenableexcept() function to enable inexactness and invalid exceptions. Once the floating-point exceptions are enabled, some computation can be performed. If some operation produces an inexact or invalid floating point number, SIGFPE is raised and a call to signal_handler is triggered. After the computation is completed, the fedisableexcept() function is called to disable the previously enabled exceptions.

But it is difficult to recover back from SIGFPE. By design, the propose usage of SIGFPE is to do some cleanup in the handler function, then exit the program gracefully. If the handler does not exit the program, after returning from the handler, the PC always points back to the instruction that caused SIGFPE and triggers SIGFPE again! However the goal is to discard the current progress and continue the program by using the LargeInteger algorithm. Although there could be potential workarounds, such as modifying the call stack and changing the return address, but implementing these solutions can be challenging, and introduces significant complexity to the codebase.

Alternatively we may read status registers and check if the imprecision bit is set:

```
bool function(matrix & tableau) {
    std::feclearexcept (FE_ALL_EXCEPT);
    if (fetestexcept(FE_INEXACT | FE_INVALID)) {
        // false for overflow, will be handle by its caller
        return false;
    } return true; // true for safe
}
```

Instead of registering floating point imprecision as interrupts, it clears the floating point status register with feclearexcept and reads it afterwards using the fetestexcept function, to check whether imprecision has ever occurred in previous computations. It then returns a boolean value to notify its caller whether or not the test for floating-point exceptions is positive, true for safe and false for overflown. In case of returning false, the caller will retry computation using LargeInteger accordingly.

3.3.3.1 Evaluation

In x86_64 there are two status registers for floating points, the legacy x87 status register for traditional scalar floating point operation, and the modern mxcsr register for SSE or AVX instructions. The source code of the library function fetestexcept indeed manipulates both registers [23]:

```
int fetestexcept(int excepts) {
    unsigned short status;
    unsigned int mxcsr;

    excepts &= FE_ALL_EXCEPT;
```

```
    /* Store the current x87 status register */
    __asm__ volatile ("fnstsw %0" : "=am" (status));

    /* Store the MXCSR register state */
    __asm__ volatile ("stmxcsr %0" : "=m" (mxcsr));

    return ((status | mxcsr) & excepts);
}
```

In the vectorized floating point scenario, the instruction for x87 status register is unnecessary and only the `mxcsr` register should be concerned. The hot loop is completely vectorized, and clang dispatches 128-bit `SSE` instructions for occasional scalar operations. This is because the `SSE` execution units can handle both single and double precision floating point arithmetic natively, whereas the legacy x87 floating-point instructions operates on an 80-bit internal format, requiring additional conversions and delay. Using `SSE` instructions reduces register pressure as well. The `SSE` units can leverage 16 `XMM` registers, but for x87 units there are only 8 floating point registers available. There could be even 32 `XMM` registers if such extension is implemented in the micro-architecture [24].

The benchmark, as illustrated in Figure 3.5, evaluates the performance of `fenv` functions alongside their respective revised versions. The modifications involves restricting operations solely to the manipulation of either the x87 status register or the `mxcsr` register. It indicates that it is significantly faster if we remove x87 status register related operations.


Figure 3.5: A benchmark for the floating point status register reading and resetting functions from `cfenv` library and their modified versions that only operates on either the x87 or the `mxcsr` status register. Excluding x87 related operations makes both `fetestexcept feclearexcept` faster.

3.3.4   Comparing `int16_t` and `float`

Section 3.3.1, 3.3.2 and 3.3.3 have concluded that `int16_t` is superior than any other integer data types and `float` is better than `double`.

Benchmark (Figure 3.3) on the toy example reveals that `int16_t`'s spends a considerably higher percentage of runtime on overflow checking, compared with the floating point data types. This is consistent with the reasoning from previous chapters, where overhead for `float` is a one-time expense, but for `int16_t` it is always a portion of the total runtime.

Note that even though `int16_t`is faster than `float`in this benchmark, it is not sound to conclude that the `pivot` function implemented in `int16_t`will be faster than `float`. The toy example is different from the `pivot`function in many aspects, for example, number of memory load operations. In case of the actual `pivot` function, float may potentially outperform int16.

Chapter 4

Implementation and Optimization of `pivot`

4.1 Matrix-wise transprecision

Transprecision techniques can be implemented at different levels of scale, such as element-wise, row-wise, and matrix-wise. The vectorized versions of `pivot` are implemented using the matrix-wise transprecision style. For `float`, the `mxcsr` register accumulates information regarding previous occurrences of overflow and imprecision, until it is cleared manually.

The element-wise method is not suitable in this scenario, as it defeats the purpose of vectorization. The row-wise method is not chosen either, due to that overflow is not likely to occur and the matrix is small. Reading the `mxcsr` can be considered as a expensive operation, comparing to the time spent on pivoting through the entire matrix. Avoiding wasted arithmetic instructions after overflow at the cost of more `mxcsr` reads is not a cost-effective trade-off.

Even though the `pivot` function using `int16_t` was implemented using the row-transprecision approach in previous works [11], it is modified to match with the matrix-wise transprecision style, in order to control differences between the `float` counterpart. After the overflow checking instructions, instead of a branch instruction pointing to the overflow handler, the boolean operator OR is applied between that overflow checking result register and a overflow flag.

4.2 Double buffering

Double buffering is a powerful technique to address the issue of data pollution caused by overflown data. When matrix-wise computing is carried out, a potential overflow can contaminate the input matrix and write meaningless results into it. It is difficult to recover from overflown results, making it impossible to dispatch the same input matrix to a algorithm of higher precision and defeating the purpose of transprecision computing.

Double buffering addresses this challenge by allocating two distinct pieces of memory, one for the input matrix and the other for output matrix. The input matrix is read-only, while the output matrix allows both reading and writing. This separation of data storage ensures that the input matrix remains unpolluted by overflown data and that any potential overflow is encapsulated within the separate output matrix.

While making a copy of the input matrix is a simple and easy solution for protecting the original data from pollution, double buffering is a superior technique due that it does not introduce additional memory operations. With double buffering, every operand requires a load operation from the input matrix, and every result takes a store operation to be written into the output matrix. This is the minimal amount of memory operation required for arithmetics. Zen4 is capable of loading one `ZMM` vector per cycle and storing one `ZMM` vector per two cycles. In comparison, it can do two multiplication or addition of `ZMM` vectors every cycle. The discrepancy in throughput between computing and IO suggests that memory copy very expensive and inefficient.

4.3 Alignment

In the Listing 3.3 and Listing 3.4, the assembly for load and store are `vmovups` and `vmovdqu64`. `vmovups` for "Move Unaligned Packed Single-Precision Floating-Point Values" [25], and `vmovdqu64` for "Move Unaligned Packed Integer Values" [26].

Unaligned load may bring potential performance impacts. When accessing memory, the CPU retrieves data from the main memory or cache in units of cache lines. On Zen4 the size of cache line is 512-bits, exactly the size of a `ZMM` register [1]. Alignment guarantees that a single cache line can cover a vector register. Otherwise, a `ZMM` register might be spitted on two cache lines. Both cache lines of 1028 bits have to be requested from the memory, then performs a shift to extract the desired 512 bits [14].

To address this issue, an aligned allocator is given to the constructor of `std::vector` when initializing a matrix [11]. The compiler is capable of recognizing alignment and issue aligned load instructions `vmovdqa` or `vmovaps` accordingly.

4.4 Reduce number of matrix index computation

Section 3.2 discussed the benefits of using a single `std::vector` to represent a matrix. Given arbitrary row and column, the item of the matrix can be found in the array by computing the index: `column count * row + column`. The matrix data structure can be further optimized if the number of index computations can be reduced. By placing the pivot row as the first row in the matrix, it is only required to compute the address once for the first element of the matrix, then every subsequent item can be indexed by incrementing the address.

4.5 Vector size specialization

Figure 2.1 shows that 95% most of the matrices have less 20 columns while 75% of the matrices have more than 18 columns [10]. A `YMM` register can pack 8 `float`s, and a `ZMM` register can pack 16 `float`s (Figure 3.4). When padding each row of the matrix to 32 columns using 2 `ZMM` registers, at least 12 `float`s are wasted. Alternatively, the rows can be padded to 24 columns using 3 `YMM` registers, and reducing the minimal padding size to 4 `float`s.

Since `YMM` registers have double amount of execution units than `ZMM` registers, while there will be only 50% more instruction count, specialize the most of the cases of less than 24 columns with 3 `YMM` registers is expected to be faster than a more genetic implementation of 32 columns and 2 `ZMM` registers.


Figure 4.1: TODO: write caption

4.6 Evaluation

Chapter 5

Conclusion and Future work

`half` is a 16-bit floating point data structure, consistent of 1-bit sign, 5-bit exponent and 10-bit mantissa (Figure 5.1). The `halfi` data type can be defined accordingly, and overflow would not occure for xx% of the cases.

Now `half` has already been supported widely on GPUs for AI applications. It provides performance benefits over higher precision floating point formats, for the exact same reason that extra precisions are a waste of resources, and more number can be computed once with a a shorter data type. If 16-bit floating point is integrated into future `AVX` extensions, it would further improve the performance of `pivot`. It is expected to be slightly slower then the `int16_t` overflow-ignored version, and the runtime can be reduced to about 50% comparing to the `pivot` function implemented using `float` in this report.


Figure 5.1: IEEE 754 half precision floating point (16 bits). [17]

Bibliography

[1] AMD. Software optimization guide for amd family 17h processors, 2021.

[2] AOIDUO, LebedevRI, RKSimon, adibiagio, and llvmbot. Un-expected rthroughput for vfmadd* instructions in znver3. https://github.com/llvm/llvm-project/issues/59325, 2022.

[3] At32Hz. Zen 2 - microarchitectures - amd, 2022.

[4] Mikolaj Bojanczyk and Joël Ouaknine. A simple and practical linear-time algorithm for presburger arithmetic. 2004.

[5] LLVM Contributors. Mlir llvm. https://mlir.llvm.org/docs/Dialects/Affine/, 2023.

[6] LLVM Contributors. Mlir llvm. https://mlir.llvm.org/, Accessed on 2023-04-03.

[7] Ian Cutress. The intel skylake-x review: Core i9 7900x, i7 7820x and i7 7800x tested, 2017.

[8] Google Benchmark Developers. google/benchmark, 2023.

[9] LLVM Developers. llvm-mca - llvm machine code analyzer, 2023.

[10] Grosser et al. Fast linear programming through transprecision computing on small and sparse data. In *Proceedings of the ACM on Programming, Languages, Volume 4*, 2020.

[11] Pitchanathan et al. Fpl: fast presburger arithmetic through transprecision. In *Proceedings of the ACM on Programming Languages, Volume 5*, 2021.

[12] Tobias Gysi, Tobias Grosser, Laurin Brandner, and Torsten Hoefler. A fast analytical model of fully associative caches. In *Proceedings of the 40th ACM SIGPLAN Conference on Programming Language Design and Implementation*. ACM, June 2019.

[13] Arch Linux. Cpu frequency scaling, 2023.

[14] Mauricio Alvarez Mesa, Esther Salamí, Alex Ramírez, and Mateo Valero. Performance impact of unaligned memory operations in simd extensions for video codec applications. DBLP, April 2007.

[15] Mysticial. Zen4's avx512 teardown, 2022.

[16] robocat. Linus torvalds on avx512, 2020.

[17] Mariana Silva, Wanjun Jiang, Matthew West, Erin Carrier, Adam Stewart, and Luke Olson. Floating point representation, 2020.

[18] TODO. Todo, TODO.

[19] TODO. Todo, TODO.

[20] TODO. Todo, TODO.

[21] TODO. Todo, TODO.

[22] TODO. Todo, TODO.

[23] TODO. Todo. TODO, TODO.

[24] TODO. Todo, TODO.

[25] TODO. Todo, TODO.

[26] TODO. Todo, TODO.

[27] Linus Torvalds. Alder lake and avx-512, 2020.